



B=8, T=0.7 | all 480 available first-turn inputs | max output 64

LayerSkip-Llama2-7B + AMD-135M



LayerSkip-Llama3-8B + Qwama-0.5B



Mean $\pm$ sample SD over 3 fresh-process seeds. *AR: 1 run, 1 GPU; speculative modes: 2 GPUs.
Input cap 512; EOS enabled; discarded tokens excluded. Width changes are measured candidates, not universal recommendations.