





llama2_b1_tree_compute3
DUET: K1/K2=6/4, exit=21 (0-based), C=3, N1/N2=12/8, U=2, W=7, trim/stream=0/0
B=1, T=0.7, 48 prompts, cap=512/96 | instrumented diagnostic; 1722 eligible steps



Wait / sync Target post Final logits Unlabelled gap Context P2 Target ready
Target pre Proxy / launch Postprocess Miss draft P1 Proxy wait Proxy received
Stage means are diagnostics, not final TPS. Proxy side-stream work is a separate lane, never added to the target stack.
Blank timeline gaps are retained; no common hit/miss compute cost is imposed. Cross-GPU times use recorded CPU/CUDA anchors.