





llama2_b8_combo_k2_1
DUET: K1/K2=3/1, exit=21 (0-based), C=3, N1/N2=6/4, U=1, W=4, trim/stream=0/0
B=8, T=0.7, 48 prompts, cap=512/96 | instrumented diagnostic; 212 eligible steps



Wait / sync Target post Final logits Unlabelled gap Context P2 Target ready
Target pre Proxy / launch Postprocess Miss draft P1 Proxy wait Proxy received

Stage means are diagnostics, not final TPS. Proxy side-stream work is a separate lane, never added to the target stack.
Blank timeline gaps are retained; no common hit/miss compute cost is imposed. Cross-GPU times use recorded CPU/CUDA anchors.