



llama2_b8_implementation_off
 DUET: K1/K2=3/2, exit=21 (0-based), C=3, N1/N2=6/4, U=1, W=4, trim/stream=0/0
 Last pass: B=8, T=0.7, N=480 | AL*=2.124, hit=76.8%, TPS*=594.11

Measured wall time by batch cache status




P1 hit Mixed hits Miss
 P2 hit Mixed hit/miss


Boundary-free steps, including slow steps



Conditional AL* and source share



All saved passes; gray = warmup



No per-phase GPU trace was recorded for this run. These are observed whole-step/source breakdowns, not inferred phase times.
 Top and lower-left panels use the last saved pass. AL* and TPS* use different boundary filters; partial batches remain in TPS*.