



llama2_b8_warm_ssd_5
SSD: K=8, F=5
Last pass: B=8, T=0.7, N=48 | AL*=2.275, hit=43.9%, TPS*=584.92


Measured wall time by batch cache status




Boundary-free steps, including slow steps



Conditional AL* and source share



All saved passes; gray = warmup



No per-phase GPU trace was recorded for this run. These are observed whole-step/source breakdowns, not inferred phase times.
Top and lower-left panels use the last saved pass. AL* and TPS* use different boundary filters; partial batches remain in TPS*.