





llama3_b1_stream_l0_profile_s1
DUET: K1/K2=4/2, exit=16 (0-based), C=3, N1/N2=8/4, U=2, W=5, trim/stream=1/1
B=1, T=0.7, 48 prompts, cap=512/128 | instrumented diagnostic; 1870 eligible steps



Wait / sync Target post Final logits Postprocess Context P2 Target ready
Prepare Proxy / launch Accept Unlabelled gap P1 Proxy wait Proxy received
Target pre Proxy compute Commit Miss draft

Stage means are diagnostics, not final TPS. Proxy side-stream work is a separate lane, never added to the target stack.
Blank timeline gaps are retained; no common hit/miss compute cost is imposed. Cross-GPU times use recorded CPU/CUDA anchors.