



llama3_b1_warm_ssd_3
SSD: K=4, F=7
Last pass: B=1, T=0.7, N=48 | AL*=2.675, hit=65.7%, TPS*=123.73


Measured wall time by batch cache status




Boundary-free steps, including slow steps



Conditional AL* and source share



All saved passes; gray = warmup



No per-phase GPU trace was recorded for this run. These are observed whole-step/source breakdowns, not inferred phase times.
Top and lower-left panels use the last saved pass. AL* and TPS* use different boundary filters; partial batches remain in TPS*.