



llama3_b1_warm_ssd_4
SSD: K=8, F=1
Last pass: B=1, T=0.7, N=48 | AL*=2.947, hit=30.2%, TPS*=95.33


Measured wall time by batch cache status




Boundary-free steps, including slow steps



Conditional AL* and source share



All saved passes; gray = warmup



No per-phase GPU trace was recorded for this run. These are observed whole-step/source breakdowns, not inferred phase times.
Top and lower-left panels use the last saved pass. AL* and TPS* use different boundary filters; partial batches remain in TPS*.