



llama3_b4_scale_ssd
SSD: K=4, F=1
Last pass: B=4, T=0.7, N=480 | AL*=2.696, hit=36.2%, TPS*=368.00


Measured wall time by batch cache status




Boundary-free steps, including slow steps



Conditional AL* and source share



All saved passes; gray = warmup



No per-phase GPU trace was recorded for this run. These are observed whole-step/source breakdowns, not inferred phase times.
Top and lower-left panels use the last saved pass. AL* and TPS* use different boundary filters; partial batches remain in TPS*.