





llama3_b8_tree_al_compute3
DUET: K1/K2=6/2, exit=30 (0-based), $\bar{C}$ =3, $\bar{N1}/\bar{N2}$ =12/4, U=2, W=7, trim/stream=0/0
B=8, T=0.7, 48 prompts, cap=512/96 | instrumented diagnostic; 75 eligible steps



Wait / sync Proxy / launch Postprocess Miss draft P1 Proxy wait Target ready
Target pre Final logits Unlabelled gap Context P2 Previous P2 Proxy received
Target post

Stage means are diagnostics, not final TPS. Proxy side-stream work is a separate lane, never added to the target stack.
Blank timeline gaps are retained; no common hit/miss compute cost is imposed. Cross-GPU times use recorded CPU/CUDA anchors.