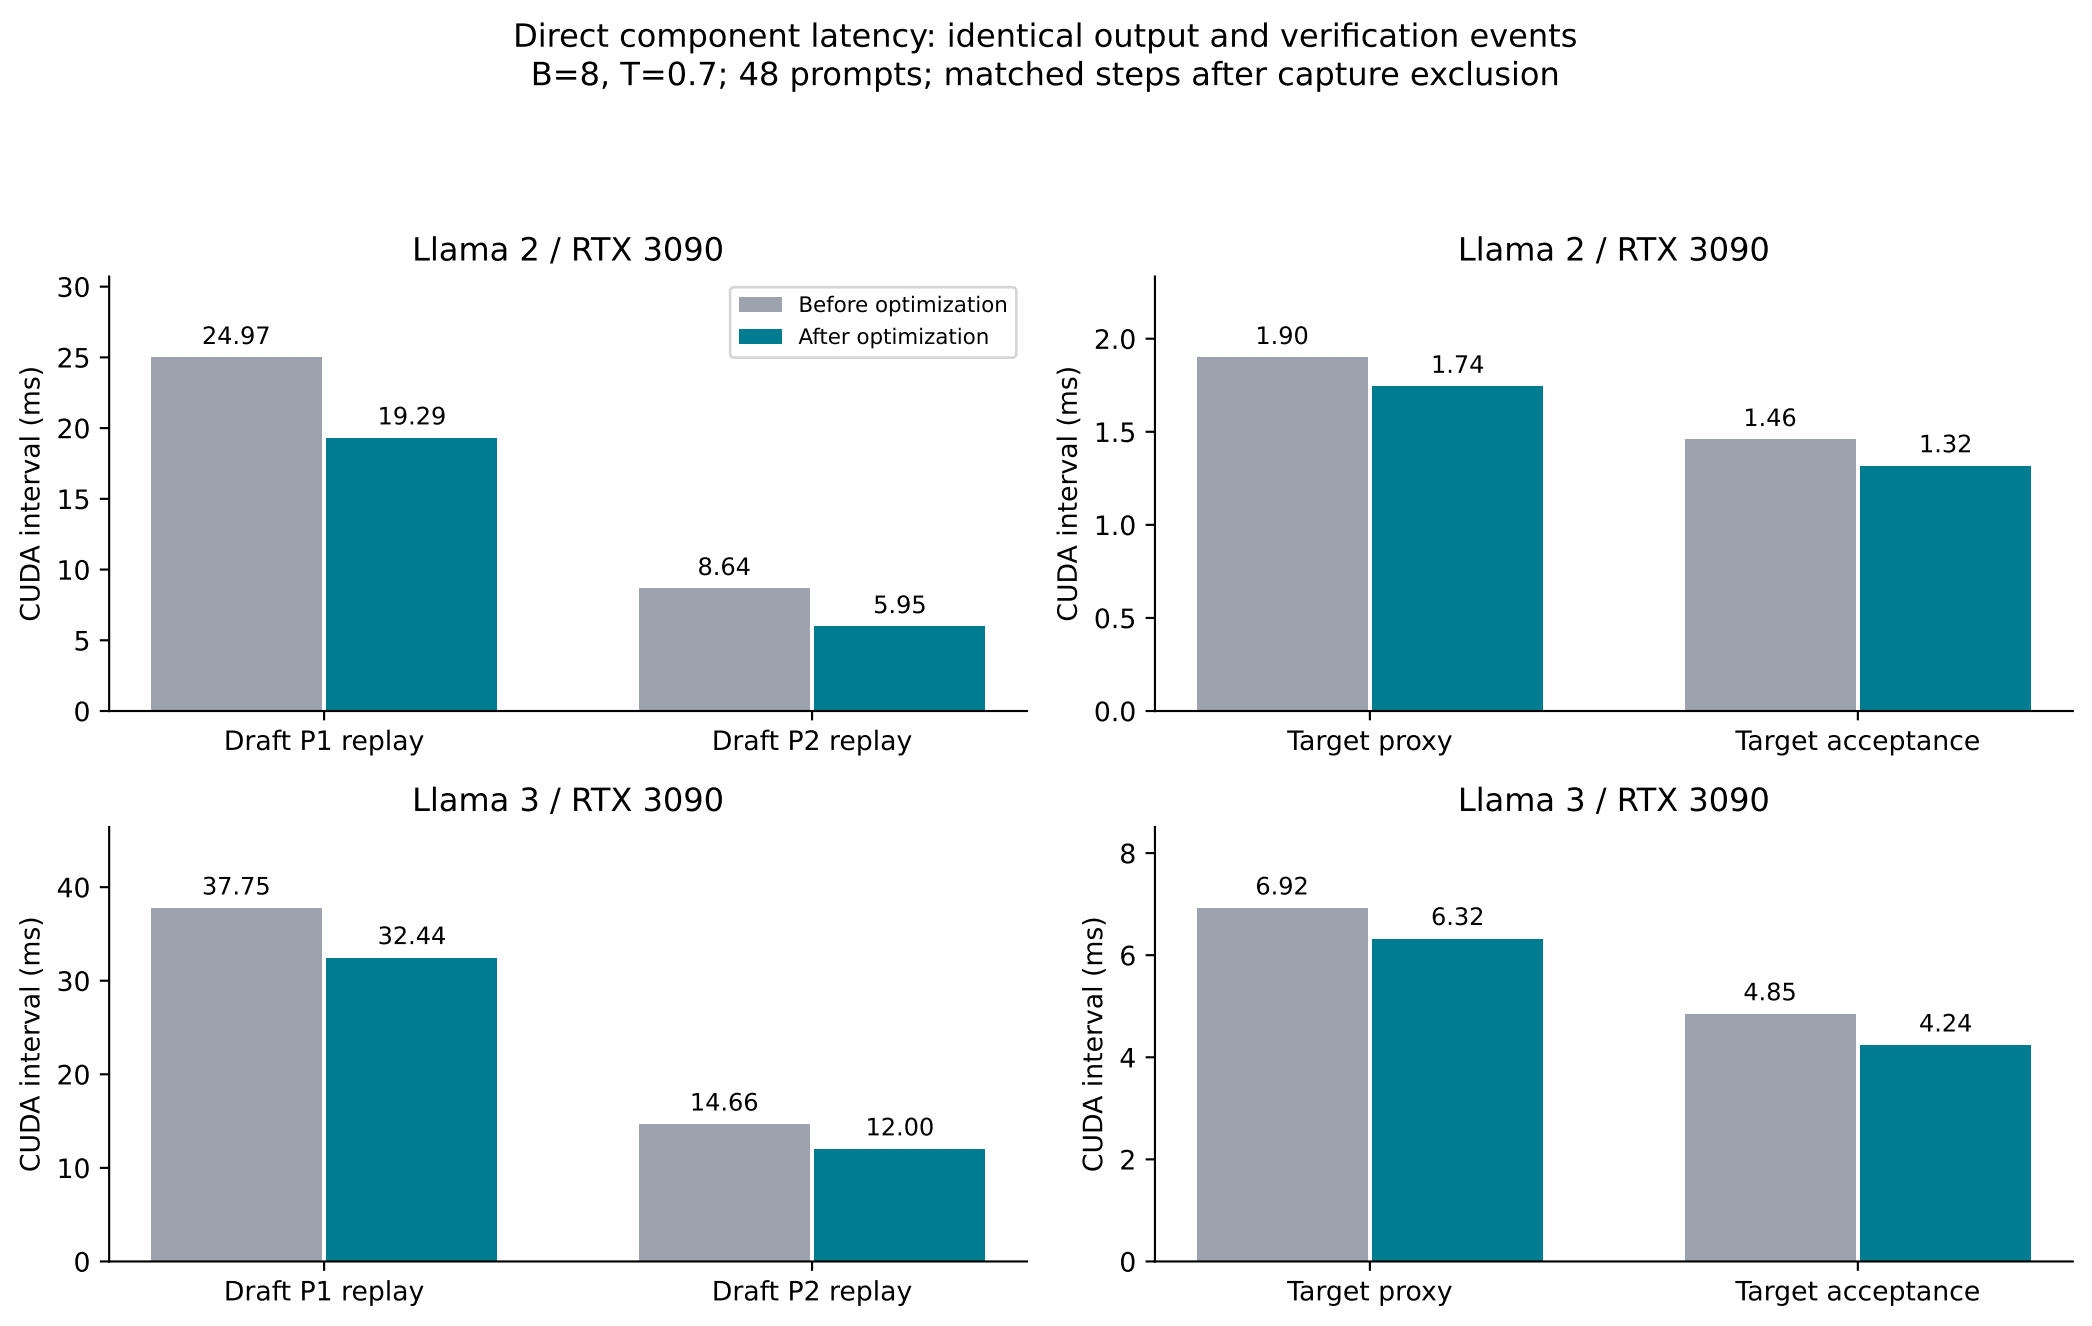
Direct component latency: identical output and verification events
B=8, T=0.7; 48 prompts; matched steps after capture exclusion
Llama 2 / RTX 3090
Before optimization
After optimization
CUDA interval (ms)
24.97
19.29
8.64
5.95
Draft P1 replay
Draft P2 replay
Llama 2 / RTX 3090
1.90
1.74
1.46
1.32
Target proxy
Target acceptance
Llama 3 / RTX 3090
37.75
32.44
14.66
12.00
Draft P1 replay
Draft P2 replay
Llama 3 / RTX 3090
6.92
6.32
4.85
4.24
Target proxy
Target acceptance