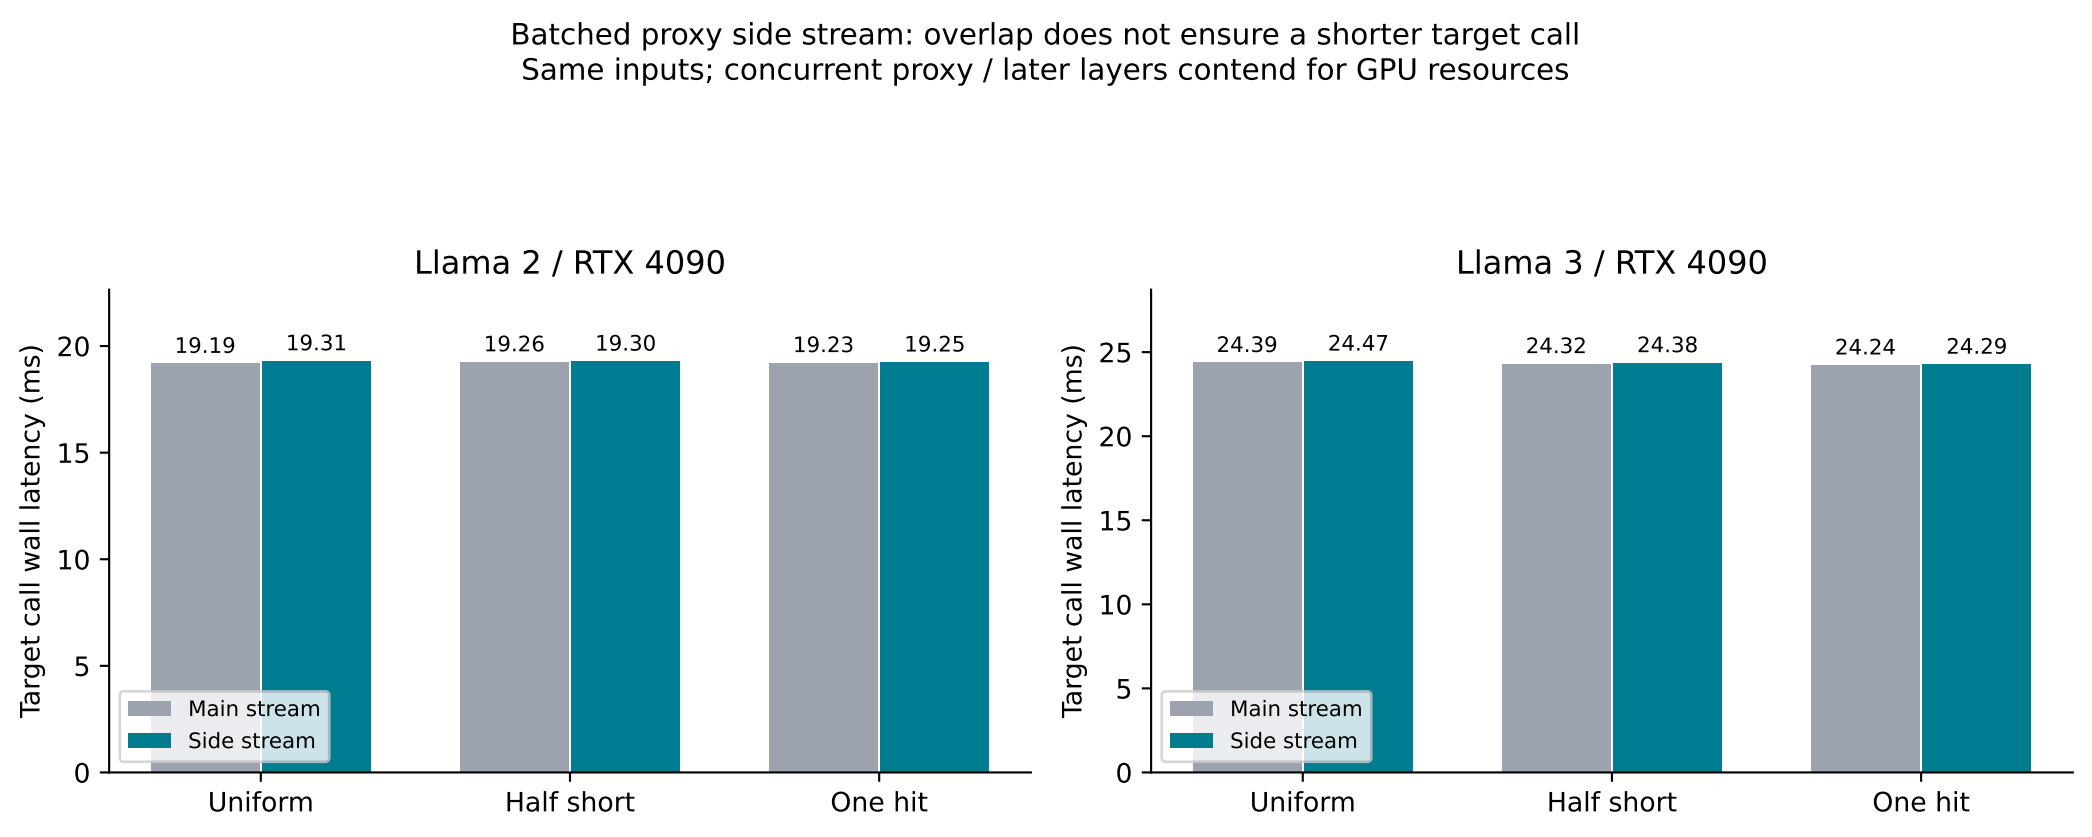
Batched proxy side stream: overlap does not ensure a shorter target call
Same inputs; concurrent proxy / later layers contend for GPU resources
Llama 2 / RTX 4090
Llama 3 / RTX 4090
Target call wall latency (ms)
Main stream
Side stream
19.19
19.31
19.26
19.30
19.23
19.25
24.39
24.47
24.32
24.38
24.24
24.29
Uniform
Half short
One hit