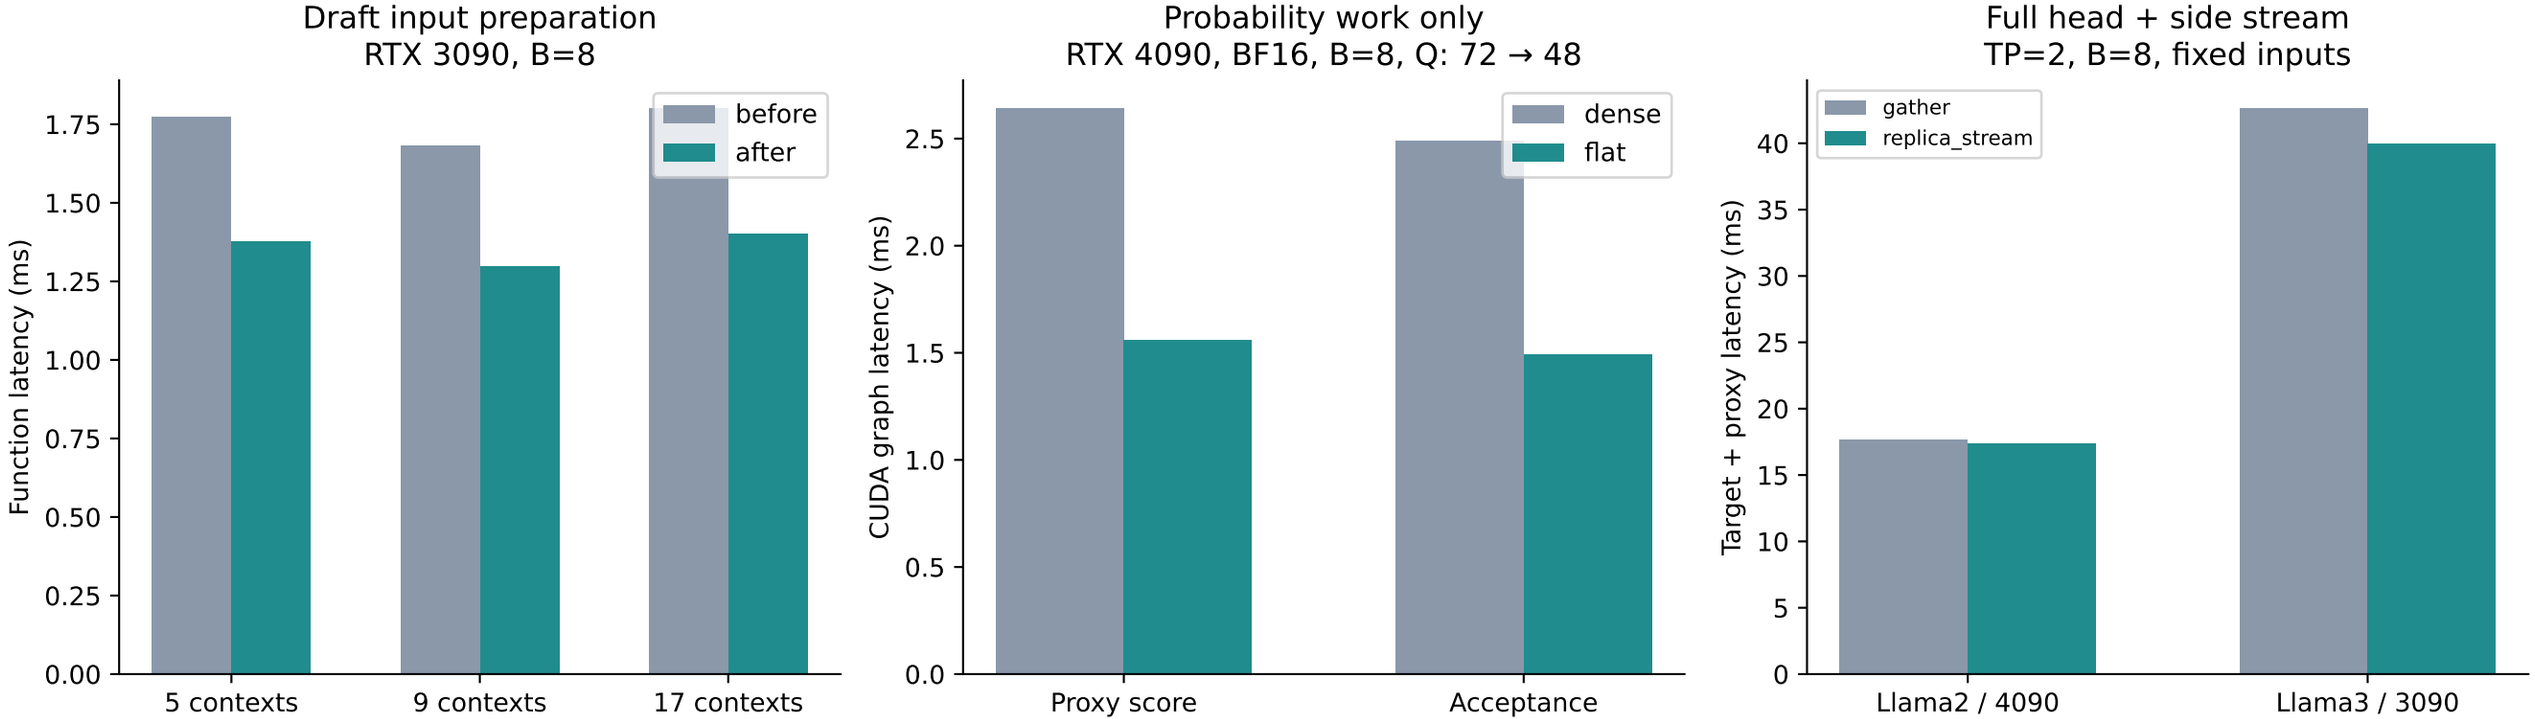
Draft input preparation
RTX 3090, B=8
Function latency (ms)
0.00
0.25
0.50
0.75
1.00
1.25
1.50
1.75
before
after
5 contexts
9 contexts
17 contexts
Probability work only
RTX 4090, BF16, B=8, Q: 72 → 48
CUDA graph latency (ms)
0.0
0.5
1.0
1.5
2.0
2.5
dense
flat
Proxy score
Acceptance
Full head + side stream
TP=2, B=8, fixed inputs
Target + proxy latency (ms)
0
5
10
15
20
25
30
35
40
gather
replica_stream
Llama2 / 4090
Llama3 / 3090