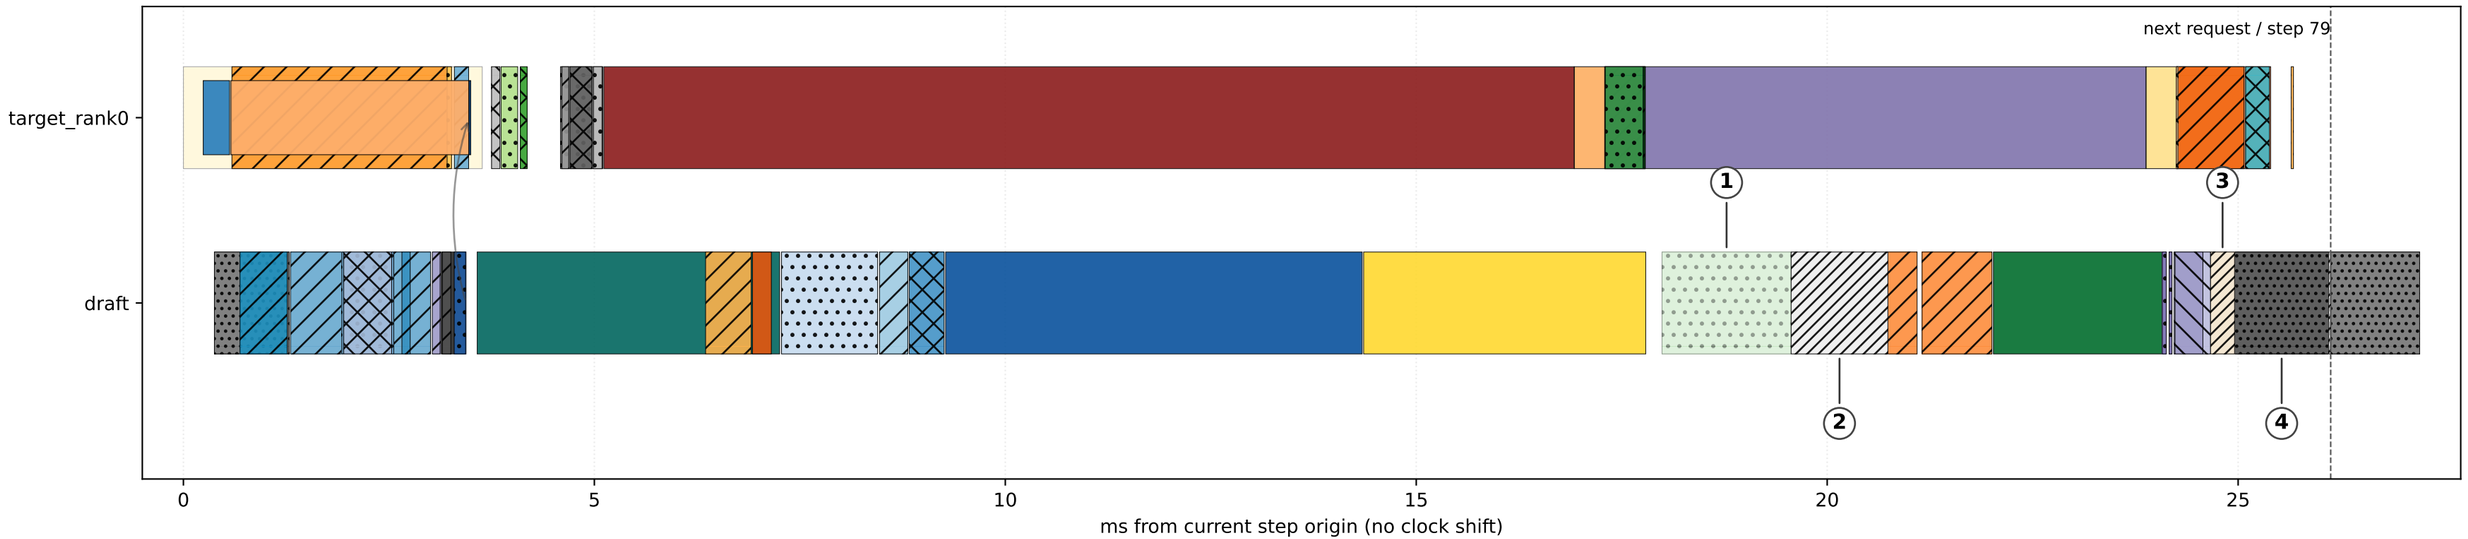
LayerSkip Llama2-7B + llama-68M
P1 hit / p50 / UID 244_t0 / request 2 / step 78

exit 20 (0-based) | K1/K2 6/3 | G=M DS/PS 12/8 | C=3 | frontier | target GPUs=1 + draft GPU=1 | B=1, BF16, T=0.7

Measured v1.0 82d964e0 | eslab18 RTX3090 | original raw-event lanes + marked gaps / next receive



1 PS roots / layout: 1.571 ms 2 Unprofiled PS dispatch: 1.177 ms (CPU gap 1.177 ms)

3 Post-PS unprofiled: 0.292 ms 4 Recorded command receive: 1.146 ms PS end -> next receive start: 1.461 ms

PS total 6.872 ms = graph 2.052 + outside graph 4.820 ms. Command -> request gap: 0.023 ms.

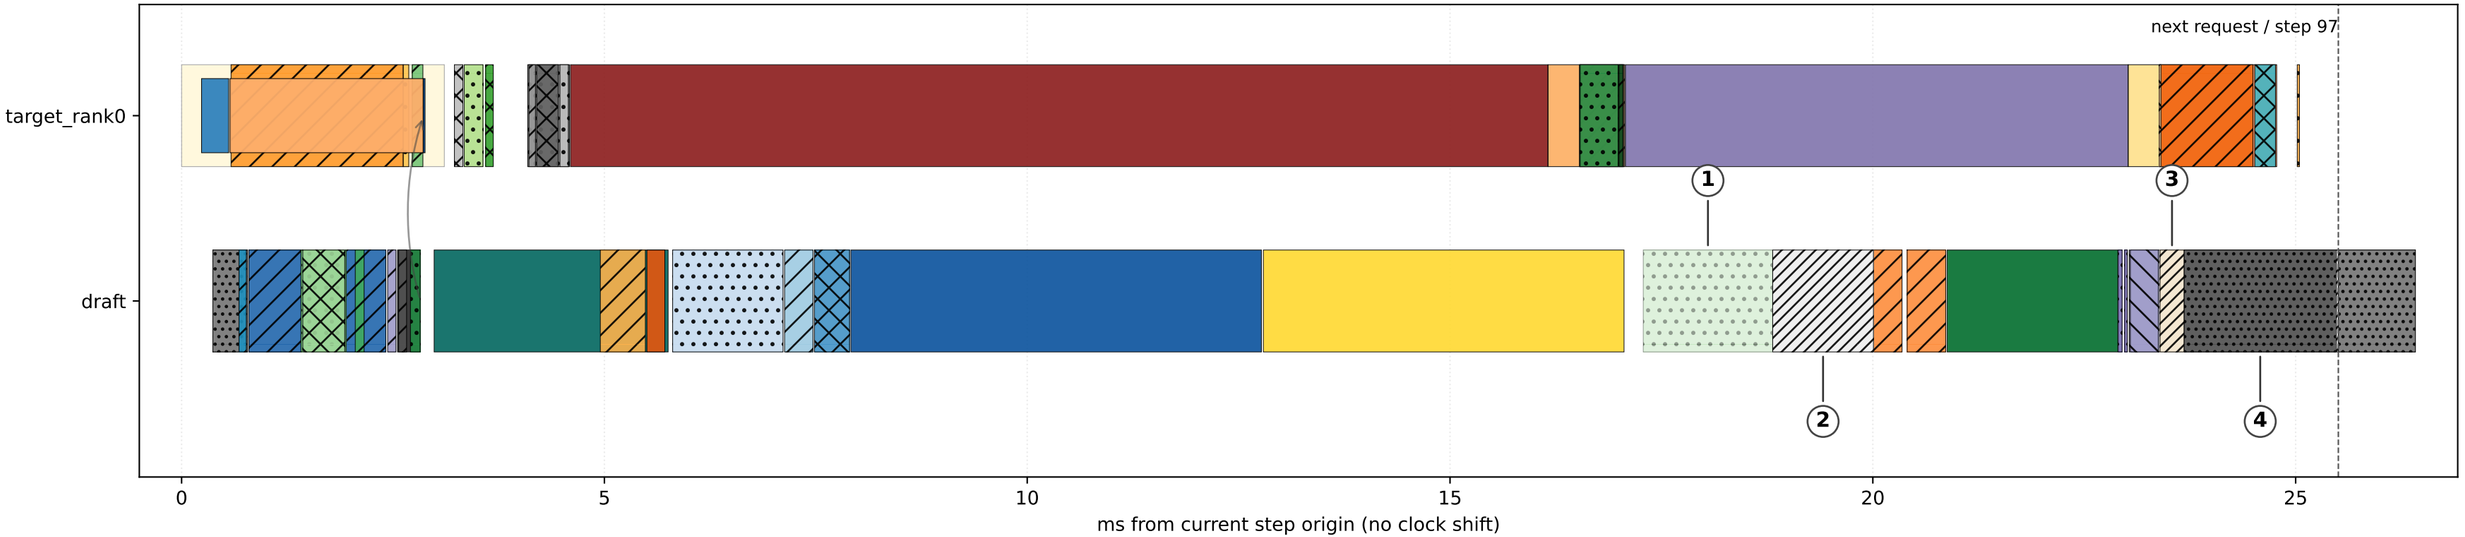
- Command-to-request gap (unprofiled)
- PS dispatch gap (unprofiled)
- Post-PS gap (unprofiled)
- draft_recv_cmd
- draft_recv_request
- draft_response_pack
- draft_send_fused
- draft_send_parent_q_p1
- draft_send_response
- exit_logits
- exit_proxy_launch
- exit_proxy_side
- final_logits
- glue
- graph_post
- graph_pre
- hit_cache_respond_hit_k1
- merge_cache
- p1_graph_replay
- p1_prepare
- p1_root_build
- p1_slot_prepare
- p2_cache_merge
- p2_graph_replay
- p2_output_convert
- p2_prepare
- phase2_build
- proxy_send_enqueue
- proxy_wait
- target_postprocess
- target_recv_fused
- target_recv_parent_q_p1
- target_recv_response_wait
- target_response_meta_read
- target_response_received
- target_send_request
- target_spec_wait
- tree_accept_walk
- tree_hit_pack_served_p1
- tree_hit_parent_q_gather_p1
- tree_hit_rerank_p1
- tree_kv_commit
- tree_kv_restore
- tree_parent_q_select
- tree_prep
- tree_proxy_graph_replay
- tree_proxy_topology_prepare
- tree_replay
- tree_verify_attention_buffers
- tree_verify_mask_prepare
- tree_verify_meta_cpu
- tree_wire_parse_validate
- verify_accept_prep
- verify_sample_accept
- verify_setup

Hatched gaps are inferred elapsed intervals, NOT proven GPU idle. Command receive includes communication/wait. Nested spans are not additive. Raw bars retain legacy 0.02 ms display floor.

LayerSkip Llama2-7B + llama-68M
P2 hit / p50 / UID 84_t0 / request 1 / step 96

exit 20 (0-based) | K1/K2 6/3 | G=M DS/PS 12/8 | C=3 | frontier | target GPUs=1 + draft GPU=1 | B=1, BF16, T=0.7

Measured v1.0 82d964e0 | eslab18 RTX3090 | original raw-event lanes + marked gaps / next receive



1 PS roots / layout: 1.531 ms 2 Unprofiled PS dispatch: 1.192 ms (CPU gap 1.192 ms)

3 Post-PS unprofiled: 0.284 ms 4 Recorded command receive: 1.803 ms PS end -> next receive start: 2.109 ms

PS total 6.339 ms = graph 2.019 + outside graph 4.320 ms. Command -> request gap: 0.022 ms.

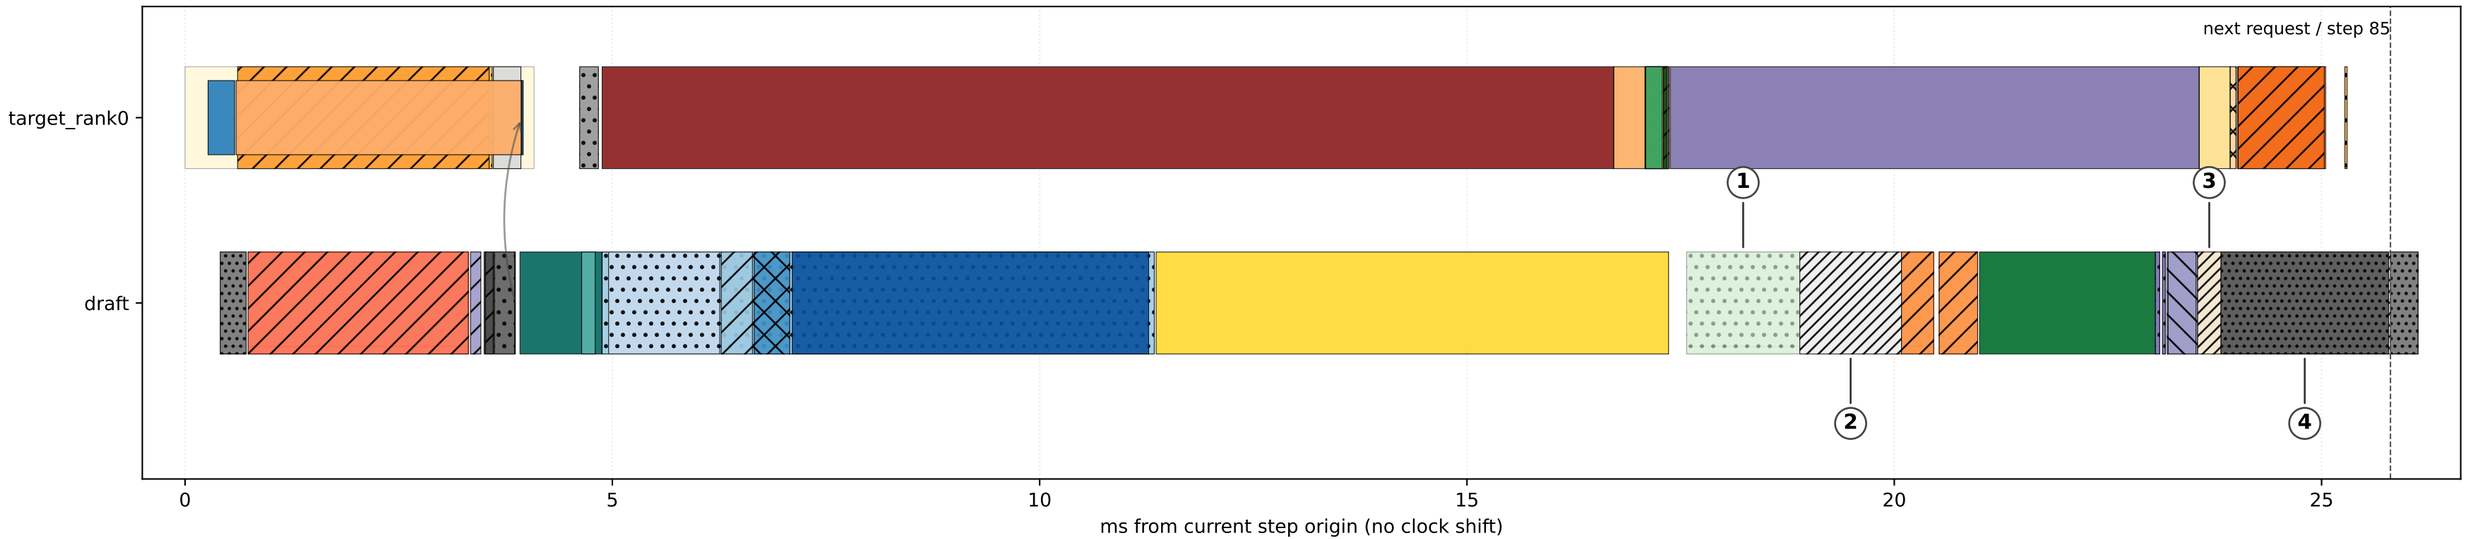
- Command-to-request gap (unprofiled)
- PS dispatch gap (unprofiled)
- Post-PS gap (unprofiled)
- draft_recv_cmd
- draft_recv_request
- draft_response_pack
- draft_send_fused
- draft_send_parent_q_p2
- draft_send_response
- exit_logits
- exit_proxy_launch
- exit_proxy_side
- final_logits
- glue
- graph_post
- graph_pre
- hit_cache_respond_hit_k2
- merge_cache
- p1_graph_replay
- p1_prepare
- p1_root_build
- p1_slot_prepare
- p2_cache_merge
- p2_graph_replay
- p2_output_convert
- p2_prepare
- phase2_build
- proxy_send_enqueue
- proxy_wait
- target_postprocess
- target_recv_fused
- target_recv_parent_q_p2
- target_recv_response_wait
- target_response_meta_read
- target_response_received
- target_send_request
- target_spec_wait
- tree_accept_walk
- tree_hit_pack_served_p2
- tree_hit_parent_q_gather_p2
- tree_hit_rerank_p2
- tree_kv_commit
- tree_kv_restore
- tree_parent_q_select
- tree_prep
- tree_proxy_graph_replay
- tree_proxy_topology_prepare
- tree_replay
- tree_verify_attention_buffers
- tree_verify_mask_prepare
- tree_verify_meta_cpu
- tree_wire_parse_validate
- verify_accept_prep
- verify_sample_accept
- verify_setup

Hatched gaps are inferred elapsed intervals, NOT proven GPU idle. Command receive includes communication/wait. Nested spans are not additive. Raw bars retain legacy 0.02 ms display floor.

LayerSkip Llama2-7B + llama-68M
Miss / p50 / UID 288_t0 / request 4 / step 84

exit 20 (0-based) | K1/K2 6/3 | G=M DS/PS 12/8 | C=3 | frontier | target GPUs=1 + draft GPU=1 | B=1, BF16, T=0.7

Measured v1.0 82d964e0 | eslab18 RTX3090 | original raw-event lanes + marked gaps / next receive



1 PS roots / layout: 1.325 ms 2 Unprofiled PS dispatch: 1.189 ms (CPU gap 1.190 ms)

3 Post-PS unprofiled: 0.271 ms 4 Recorded command receive: 1.963 ms PS end -> next receive start: 2.256 ms

PS total 6.191 ms = graph 2.056 + outside graph 4.136 ms. Command -> request gap: 0.021 ms.

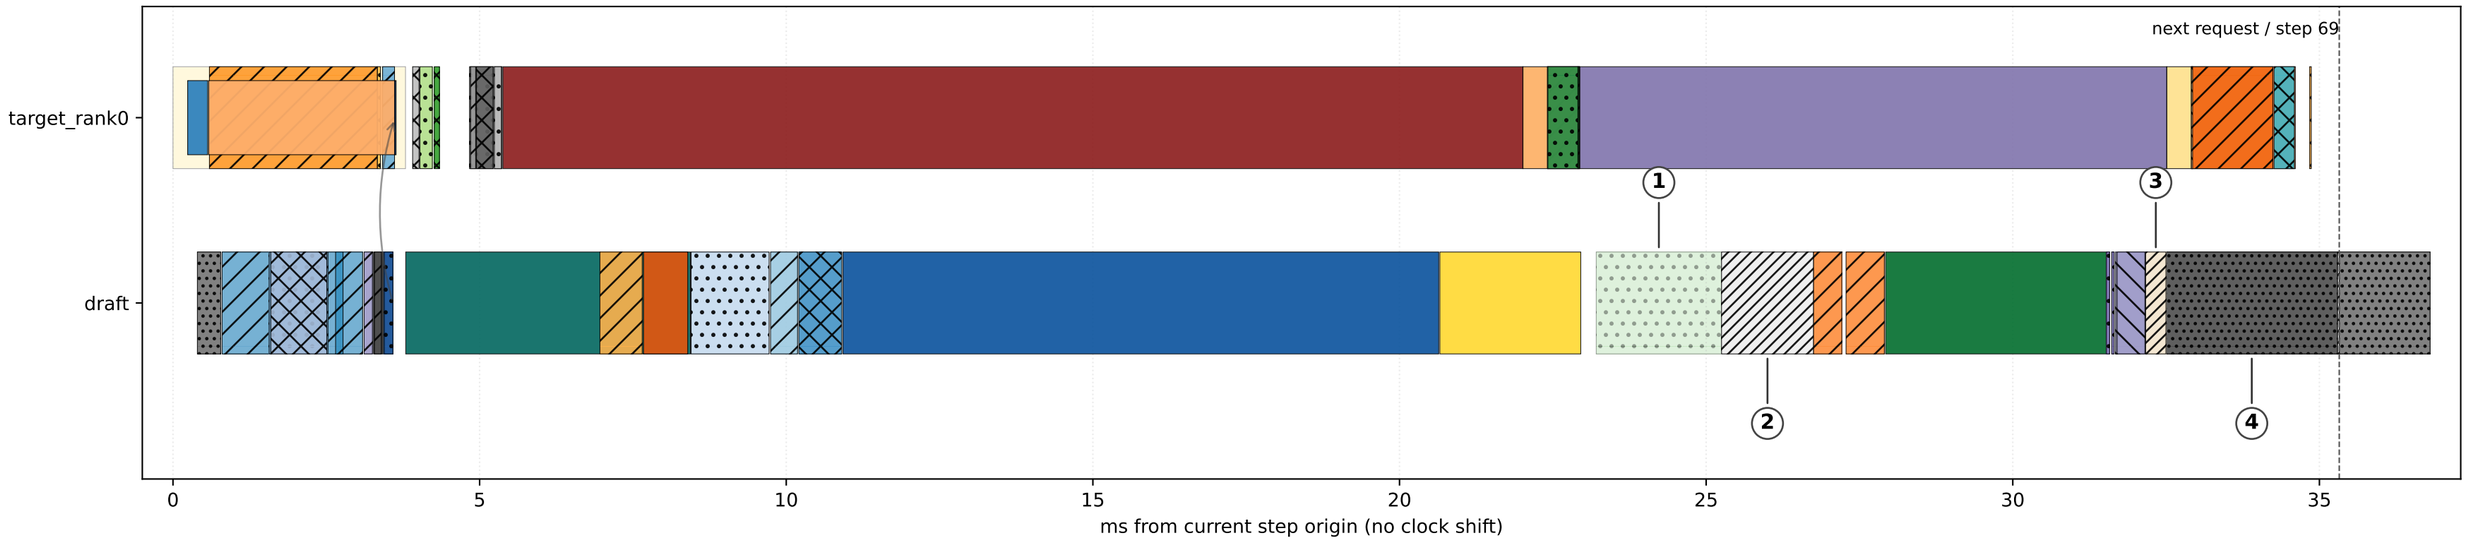
- Command-to-request gap (unprofiled)
- PS dispatch gap (unprofiled)
- Post-PS gap (unprofiled)
- chain_accept
- chain_proxy_graph_replay
- draft_glue_replay
- draft_recv_cmd
- draft_recv_request
- draft_response_pack
- draft_send_chain_q
- draft_send_fused
- draft_send_response
- exit_logits
- exit_proxy_launch
- exit_proxy_side
- final_logits
- glue
- graph_post
- graph_pre
- hit_cache_respond_miss
- merge_cache
- p1_graph_replay
- p1_prepare
- p1_root_build
- p1_slot_prepare
- p2_cache_merge
- p2_graph_replay
- p2_output_convert
- p2_prepare
- phase1_build
- phase2_build
- proxy_send_enqueue
- proxy_wait
- target_postprocess
- target_recv_chain_q
- target_recv_fused
- target_recv_response_wait
- target_response_meta_read
- target_response_received
- target_send_request
- target_spec_wait
- verify_accept_prep
- verify_sample_accept
- verify_setup

Hatched gaps are inferred elapsed intervals, NOT proven GPU idle. Command receive includes communication/wait. Nested spans are not additive. Raw bars retain legacy 0.02 ms display floor.

LayerSkip Llama2-13B + AMD-Llama-135M
P1 hit / p50 / UID 84_t0 / request 1 / step 68

exit 23 (0-based) | K1/K2 6/3 | G=M DS/PS 12/8 | C=3 | last | target GPUs=2 + draft GPU=1 | B=1, BF16, T=0.7

Measured v1.0 59f46511 | eslab18 RTX3090 | original raw-event lanes + marked gaps / next receive



1 PS roots / layout: 2.043 ms 2 Unprofiled PS dispatch: 1.499 ms (CPU gap 1.498 ms)

3 Post-PS unprofiled: 0.330 ms 4 Recorded command receive: 2.799 ms PS end -> next receive start: 3.156 ms

PS total 9.212 ms = graph 3.598 + outside graph 5.614 ms. Command -> request gap: 0.027 ms.

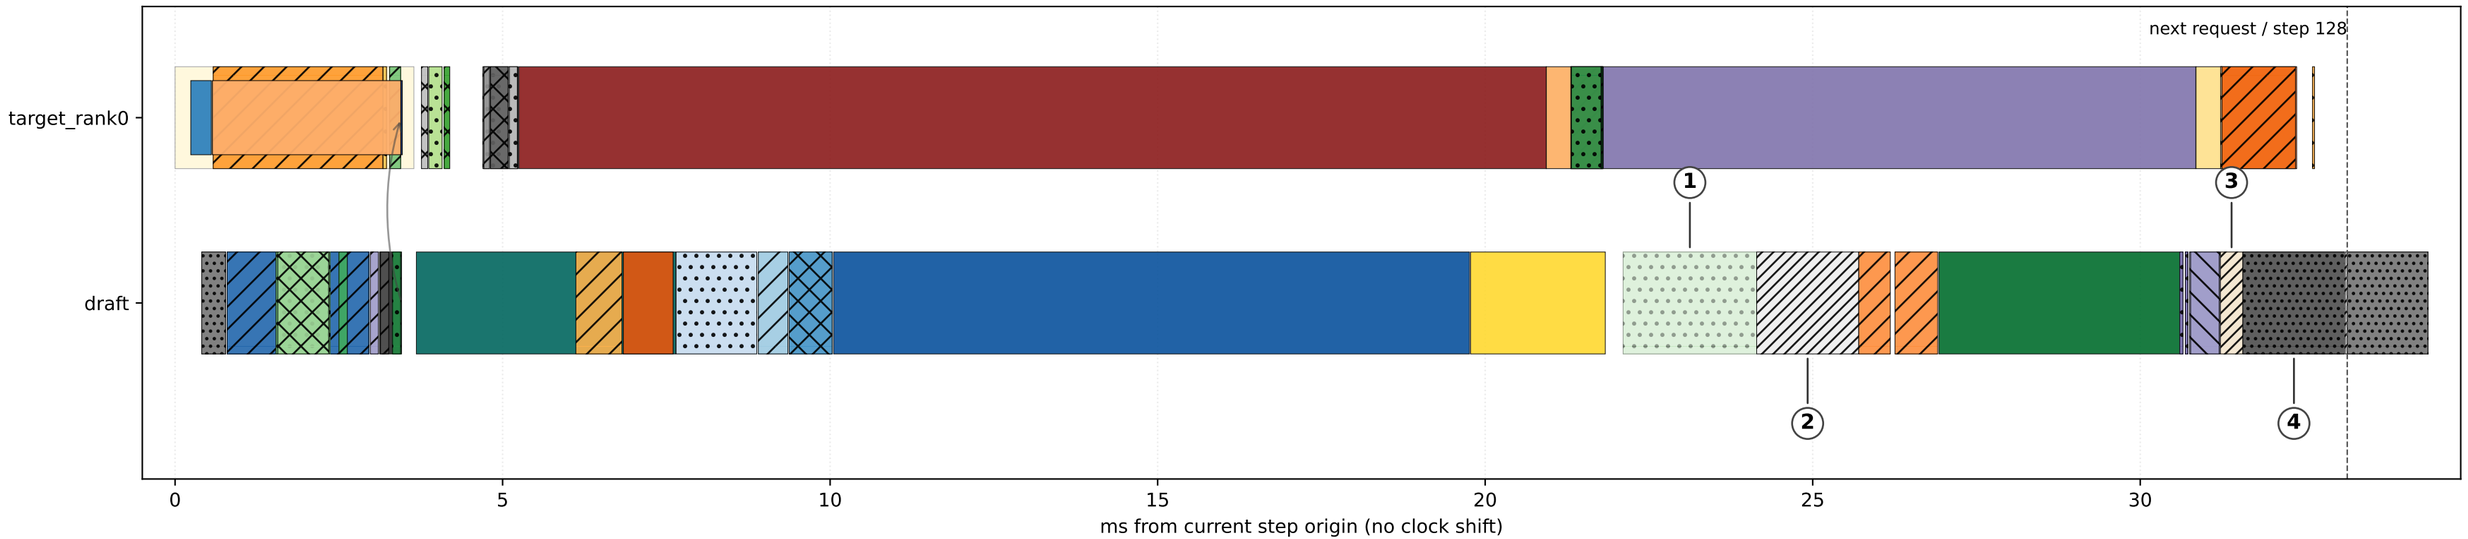
Command-to-request gap (unprofiled)	exit_logits	p1_graph_replay	proxy_send_enqueue	target_spec_wait	tree_proxy_topology_prepare
PS dispatch gap (unprofiled)	exit_proxy_launch	p1_prepare	proxy_wait	tree_accept_walk	tree_replay
Post-PS gap (unprofiled)	exit_proxy_side	p1_root_build	target_postprocess	tree_hit_pack_served_p1	tree_verify_attention_buffers
draft_recv_cmd	final_logits	p1_slot_prepare	target_recv_fused	tree_hit_parent_q_gather_p1	tree_verify_mask_prepare
draft_recv_request	glue	p2_cache_merge	target_recv_parent_q_p1	tree_hit_rerank_p1	tree_verify_meta_cpu
draft_response_pack	graph_post	p2_graph_replay	target_recv_response_wait	tree_kv_commit	tree_wire_parse_validate
draft_send_fused	graph_pre	p2_output_convert	target_response_meta_read	tree_parent_q_select	verify_accept_prep
draft_send_parent_q_p1	hit_cache_respond_hit_k1	p2_prepare	target_response_received	tree_prep	verify_sample_accept
draft_send_response	merge_cache	phase2_build	target_send_request	tree_proxy_graph_replay	verify_setup

Hatched gaps are inferred elapsed intervals, NOT proven GPU idle. Command receive includes communication/wait. Nested spans are not additive. Raw bars retain legacy 0.02 ms display floor.

LayerSkip Llama2-13B + AMD-Llama-135M
P2 hit / p50 / UID 288_t0 / request 4 / step 127

exit 23 (0-based) | K1/K2 6/3 | G=M DS/PS 12/8 | C=3 | last | target GPUs=2 + draft GPU=1 | B=1, BF16, T=0.7

Measured v1.0 59f46511 | eslab18 RTX3090 | original raw-event lanes + marked gaps / next receive



1 PS roots / layout: 2.039 ms 2 Unprofiled PS dispatch: 1.557 ms (CPU gap 1.559 ms)

3 Post-PS unprofiled: 0.340 ms 4 Recorded command receive: 1.564 ms PS end -> next receive start: 1.936 ms

PS total 9.391 ms = graph 3.674 + outside graph 5.717 ms. Command -> request gap: 0.031 ms.

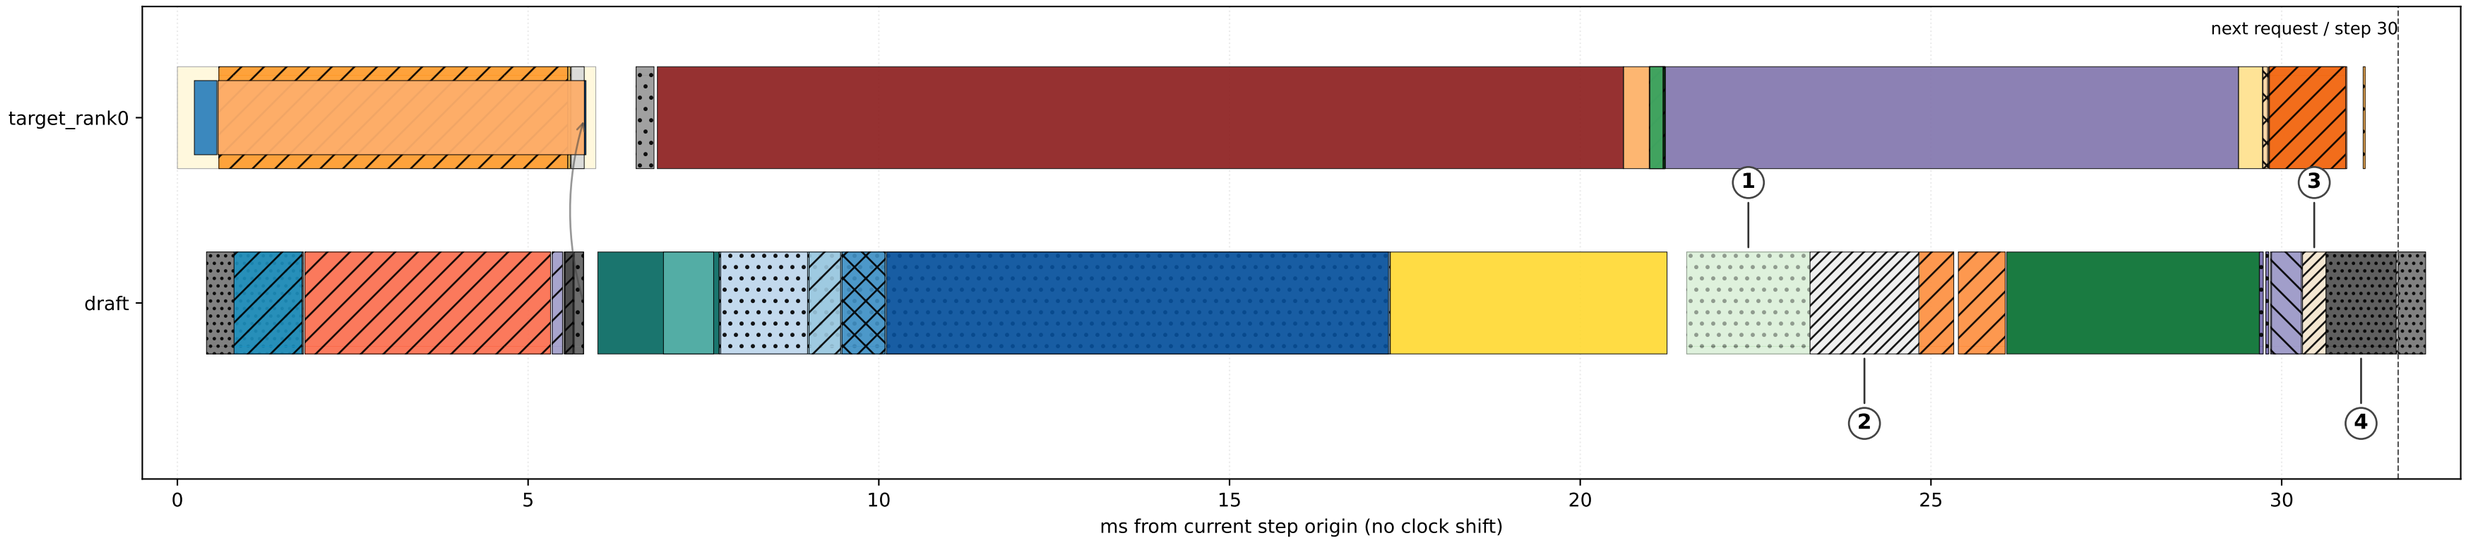
Command-to-request gap (unprofiled)
PS dispatch gap (unprofiled)
Post-PS gap (unprofiled)
draft_recv_cmd
draft_recv_request
draft_response_pack
draft_send_fused
draft_send_parent_q_p2
draft_send_response
exit_logits
exit_proxy_launch
exit_proxy_side
final_logits
glue
graph_post
graph_pre
hit_cache_respond_hit_k2
merge_cache
p1_graph_replay
p1_prepare
p1_root_build
p1_slot_prepare
p2_cache_merge
p2_graph_replay
p2_output_convert
p2_prepare
phase2_build
proxy_send_enqueue
proxy_wait
target_postprocess
target_recv_fused
target_recv_parent_q_p2
target_recv_response_wait
target_response_meta_read
target_response_received
target_send_request
target_spec_wait
tree_accept_walk
tree_hit_pack_served_p2
tree_hit_parent_q_gather_p2
tree_hit_rerank_p2
tree_parent_q_select
tree_prep
tree_proxy_graph_replay
tree_proxy_topology_prepare
tree_replay
tree_verify_attention_buffers
tree_verify_mask_prepare
tree_verify_meta_cpu
tree_wire_parse_validate
verify_accept_prep
verify_sample_accept
verify_setup

Hatched gaps are inferred elapsed intervals, NOT proven GPU idle. Command receive includes communication/wait. Nested spans are not additive. Raw bars retain legacy 0.02 ms display floor.

LayerSkip Llama2-13B + AMD-Llama-135M
Miss / p50 / UID 128_t0 / request 3 / step 29

exit 23 (0-based) | K1/K2 6/3 | G=M DS/PS 12/8 | C=3 | last | target GPUs=2 + draft GPU=1 | B=1, BF16, T=0.7

Measured v1.0 59f46511 | eslab18 RTX3090 | original raw-event lanes + marked gaps / next receive



1 PS roots / layout: 1.759 ms 2 Unprofiled PS dispatch: 1.551 ms (CPU gap 1.550 ms)

3 Post-PS unprofiled: 0.333 ms 4 Recorded command receive: 1.003 ms PS end -> next receive start: 1.363 ms

PS total 9.061 ms = graph 3.603 + outside graph 5.458 ms. Command -> request gap: 0.027 ms.

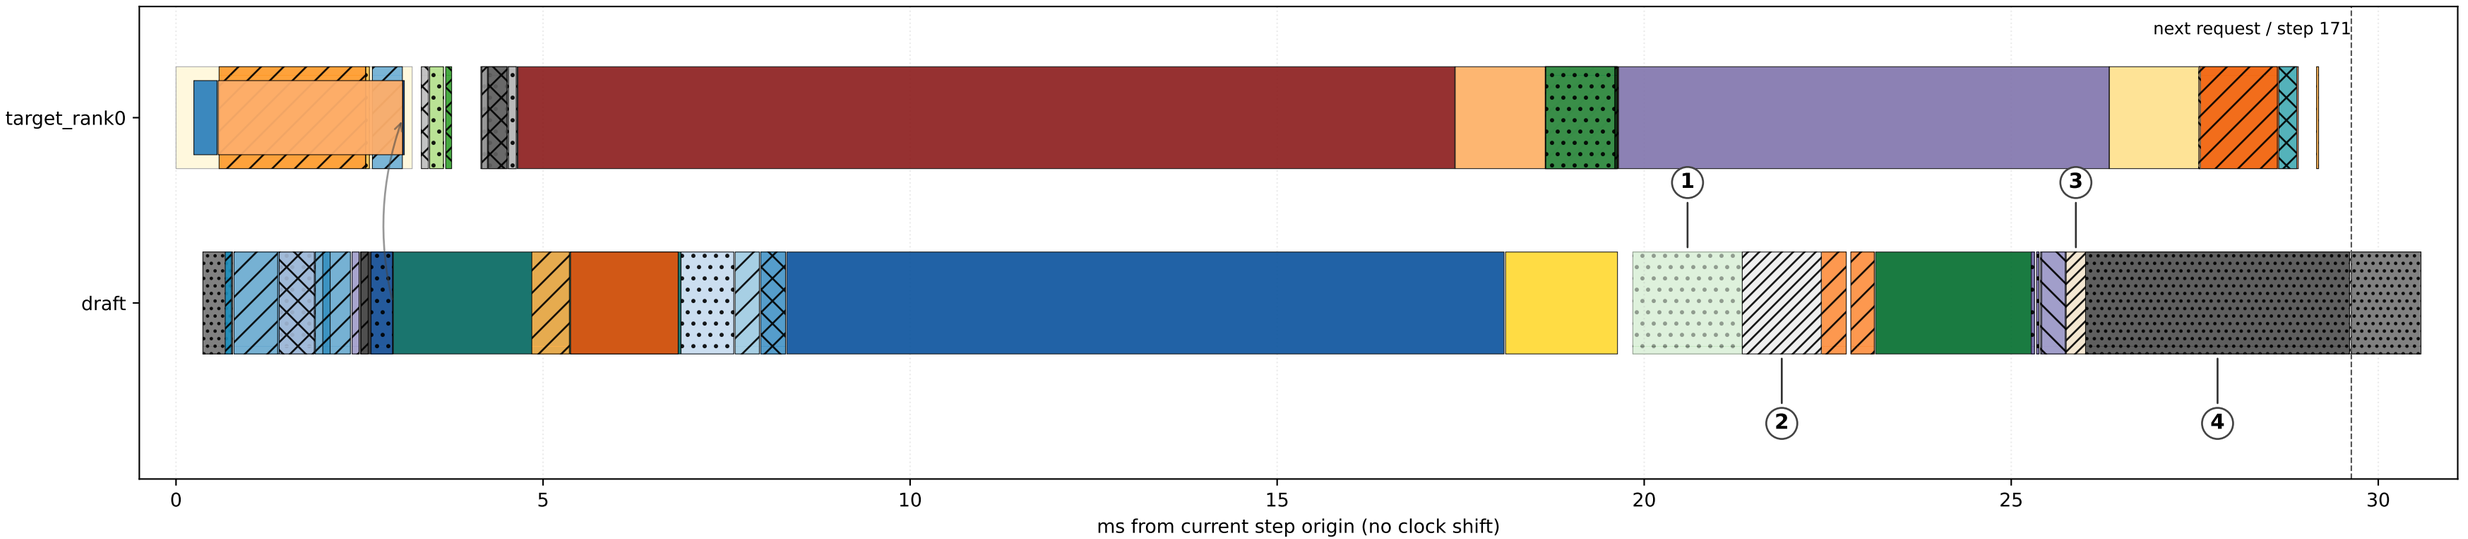
Command-to-request gap (unprofiled)
PS dispatch gap (unprofiled)
Post-PS gap (unprofiled)
chain_accept
chain_proxy_graph_replay
draft_glue_replay
draft_recv_cmd
draft_recv_request
draft_response_pack
draft_send_chain_q
draft_send_fused
draft_send_response
exit_logits
exit_proxy_launch
exit_proxy_side
final_logits
glue
graph_post
graph_pre
hit_cache_respond_miss
merge_cache
p1_graph_replay
p1_prepare
p1_root_build
p1_slot_prepare
p2_cache_merge
p2_graph_replay
p2_output_convert
p2_prepare
phase1_build
phase2_build
proxy_send_enqueue
proxy_wait
target_postprocess
target_recv_fused
target_recv_response_wait
target_response_meta_read
target_response_received
target_send_request
target_spec_wait
tree_kv_restore
verify_accept_prep
verify_sample_accept
verify_setup

Hatched gaps are inferred elapsed intervals, NOT proven GPU idle. Command receive includes communication/wait. Nested spans are not additive. Raw bars retain legacy 0.02 ms display floor.

LayerSkip Llama3-8B + Axcel Llama3.1 drafter-340M
P1 hit / p50 / UID 84_t0 / request 1 / step 170

exit 20 (0-based) | K1/K2 4/1 | G=M DS/PS 8/4 | C=3 | last | target GPUs=1 + draft GPU=1 | B=1, BF16, T=0.7

Measured v1.0 16dd4198 | eslab18 RTX3090 | original raw-event lanes + marked gaps / next receive



1 PS roots / layout: 1.491 ms 2 Unprofiled PS dispatch: 1.076 ms (CPU gap 1.076 ms)

3 Post-PS unprofiled: 0.263 ms 4 Recorded command receive: 3.598 ms PS end -> next receive start: 3.884 ms

PS total 6.113 ms = graph 2.120 + outside graph 3.993 ms. Command -> request gap: 0.023 ms.

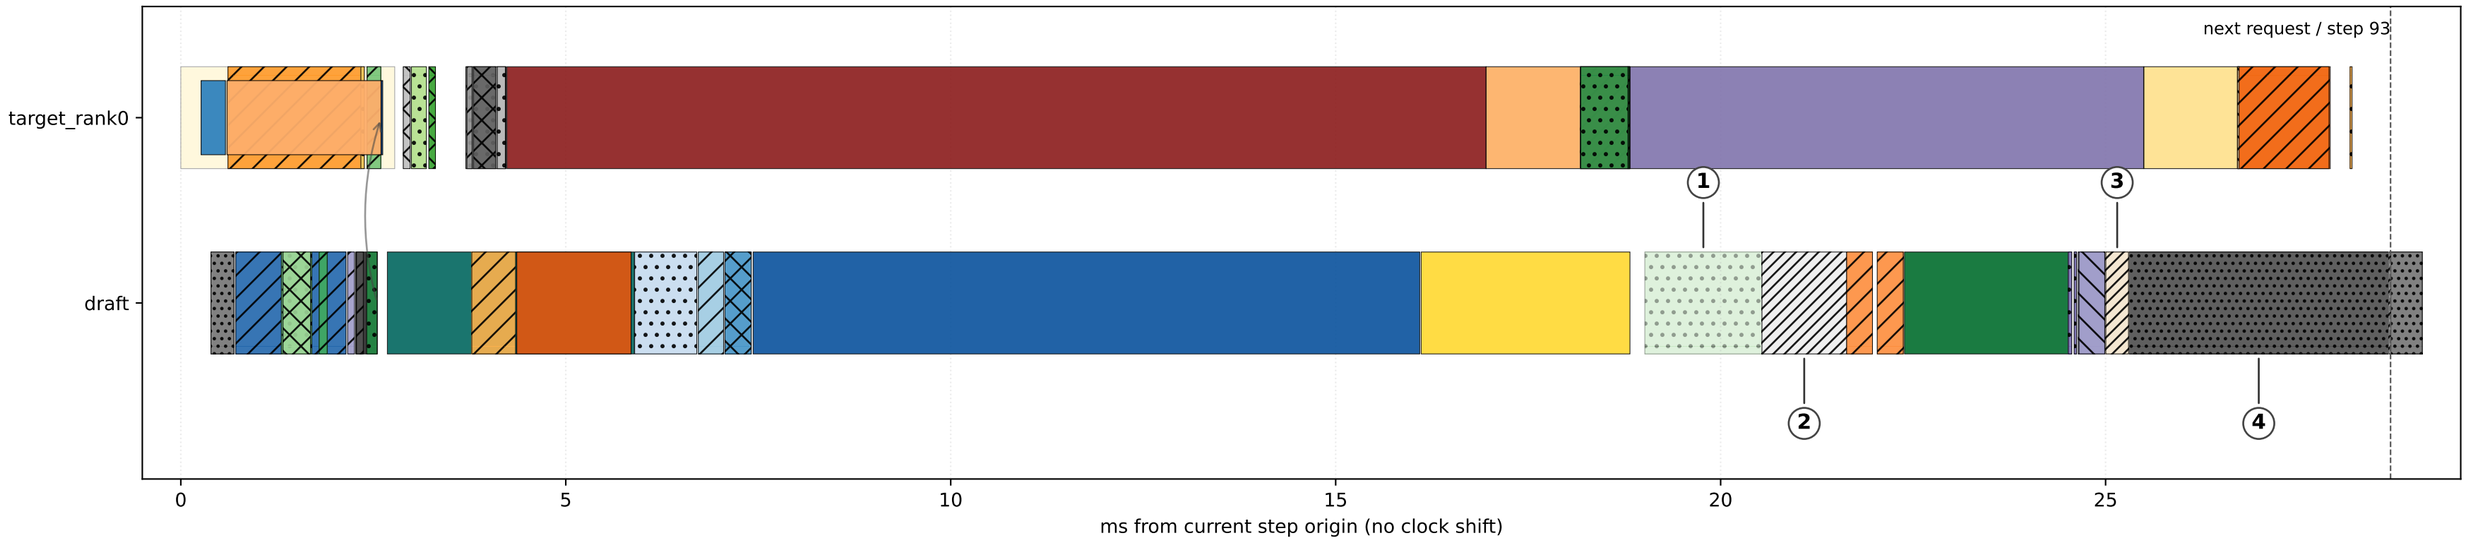
Command-to-request gap (unprofiled)
PS dispatch gap (unprofiled)
Post-PS gap (unprofiled)
draft_recv_cmd
draft_recv_request
draft_response_pack
draft_send_fused
draft_send_parent_q_p1
draft_send_response
exit_logits
exit_proxy_launch
exit_proxy_side
final_logits
glue
graph_post
graph_pre
hit_cache_respond_hit_k1
merge_cache
p1_graph_replay
p1_prepare
p1_root_build
p1_slot_prepare
p2_cache_merge
p2_graph_replay
p2_output_convert
p2_prepare
phase2_build
proxy_send_enqueue
proxy_wait
target_postprocess
target_recv_fused
target_recv_parent_q_p1
target_recv_response_wait
target_response_meta_read
target_response_received
target_send_request
target_spec_wait
tree_accept_walk
tree_hit_pack_served_p1
tree_hit_parent_q_gather_p1
tree_hit_rerank_p1
tree_kv_commit
tree_kv_restore
tree_parent_q_select
tree_prep
tree_proxy_graph_replay
tree_proxy_topology_prepare
tree_replay
tree_verify_attention_buffers
tree_verify_mask_prepare
tree_verify_meta_cpu
tree_wire_parse_validate
verify_accept_prep
verify_sample_accept
verify_setup

Hatched gaps are inferred elapsed intervals, NOT proven GPU idle. Command receive includes communication/wait. Nested spans are not additive. Raw bars retain legacy 0.02 ms display floor.

LayerSkip Llama3-8B + Axcel Llama3.1 drafter-340M
P2 hit / p50 / UID 244_t0 / request 2 / step 92

exit 20 (0-based) | K1/K2 4/1 | G=M DS/PS 8/4 | C=3 | last | target GPUs=1 + draft GPU=1 | B=1, BF16, T=0.7

Measured v1.0 16dd4198 | eslab18 RTX3090 | original raw-event lanes + marked gaps / next receive



1 PS roots / layout: 1.521 ms 2 Unprofiled PS dispatch: 1.099 ms (CPU gap 1.098 ms)

3 Post-PS unprofiled: 0.295 ms 4 Recorded command receive: 3.383 ms PS end -> next receive start: 3.698 ms

PS total 6.181 ms = graph 2.125 + outside graph 4.056 ms. Command -> request gap: 0.021 ms.

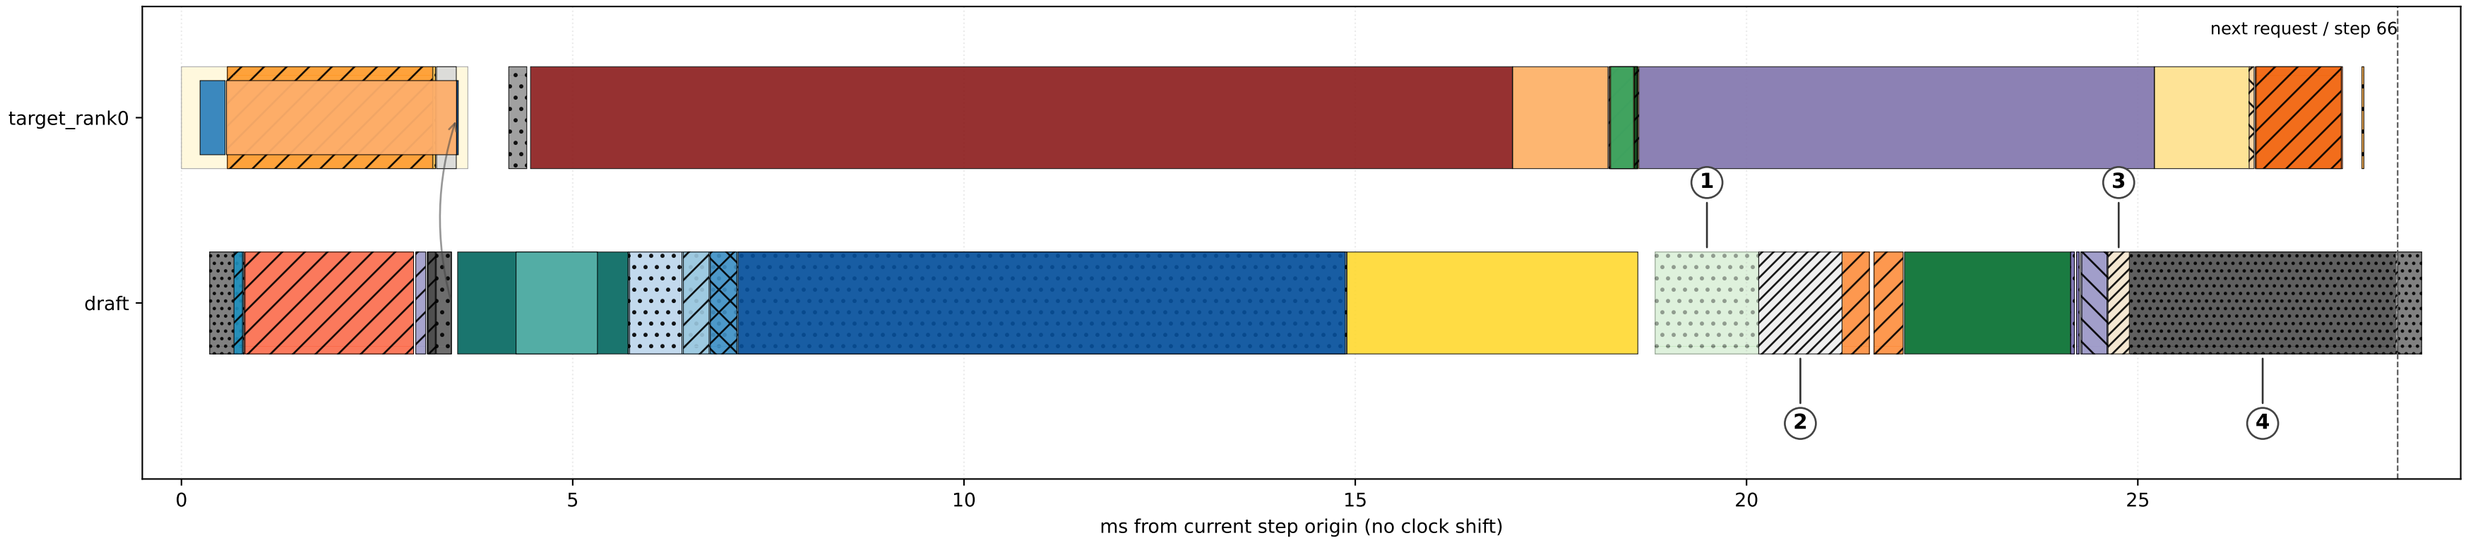
Command-to-request gap (unprofiled)	exit_logits	p1_graph_replay	proxy_send_enqueue	target_spec_wait	tree_replay
PS dispatch gap (unprofiled)	exit_proxy_launch	p1_prepare	proxy_wait	tree_accept_walk	tree_verify_attention_buffers
Post-PS gap (unprofiled)	exit_proxy_side	p1_root_build	target_postprocess	tree_hit_pack_served_p2	tree_verify_mask_prepare
draft_recv_cmd	final_logits	p1_slot_prepare	target_recv_fused	tree_hit_parent_q_gather_p2	tree_verify_meta_cpu
draft_recv_request	glue	p2_cache_merge	target_recv_parent_q_p2	tree_hit_rerank_p2	tree_wire_parse_validate
draft_response_pack	graph_post	p2_graph_replay	target_recv_response_wait	tree_parent_q_select	verify_accept_prep
draft_send_fused	graph_pre	p2_output_convert	target_response_meta_read	tree_prep	verify_sample_accept
draft_send_parent_q_p2	hit_cache_respond_hit_k2	p2_prepare	target_response_received	tree_proxy_graph_replay	verify_setup
draft_send_response	merge_cache	phase2_build	target_send_request	tree_proxy_topology_prepare	

Hatched gaps are inferred elapsed intervals, NOT proven GPU idle. Command receive includes communication/wait. Nested spans are not additive. Raw bars retain legacy 0.02 ms display floor.

LayerSkip Llama3-8B + Axcel Llama3.1 drafter-340M
Miss / p50 / UID 244_t0 / request 2 / step 65

exit 20 (0-based) | K1/K2 4/1 | G=M DS/PS 8/4 | C=3 | last | target GPUs=1 + draft GPU=1 | B=1, BF16, T=0.7

Measured v1.0 16dd4198 | eslab18 RTX3090 | original raw-event lanes + marked gaps / next receive



1 PS roots / layout: 1.323 ms 2 Unprofiled PS dispatch: 1.064 ms (CPU gap 1.075 ms)

3 Post-PS unprofiled: 0.272 ms 4 Recorded command receive: 3.407 ms PS end -> next receive start: 3.701 ms

PS total 6.009 ms = graph 2.121 + outside graph 3.888 ms. Command -> request gap: 0.021 ms.

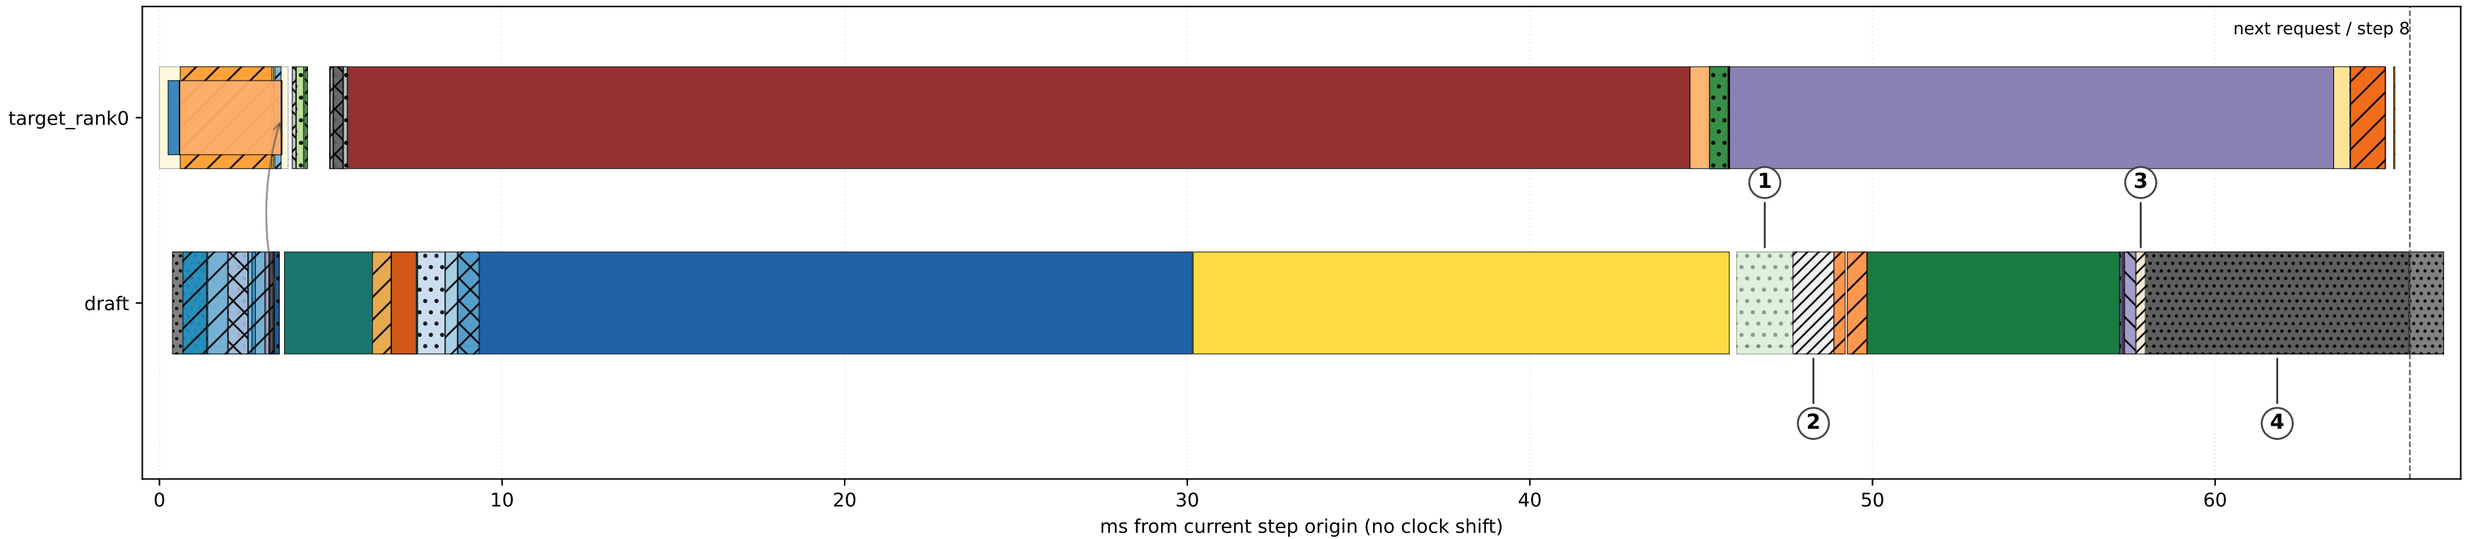
- Command-to-request gap (unprofiled)
- PS dispatch gap (unprofiled)
- Post-PS gap (unprofiled)
- chain_accept
- chain_proxy_graph_replay
- draft_glue_replay
- draft_recv_cmd
- draft_recv_request
- draft_response_pack
- draft_send_chain_q
- draft_send_fused
- draft_send_response
- exit_logits
- exit_proxy_launch
- exit_proxy_side
- final_logits
- glue
- graph_post
- graph_pre
- hit_cache_respond_miss
- merge_cache
- p1_graph_replay
- p1_prepare
- p1_root_build
- p1_slot_prepare
- p2_cache_merge
- p2_graph_replay
- p2_output_convert
- p2_prepare
- phase1_build
- phase2_build
- proxy_send_enqueue
- proxy_wait
- target_postprocess
- target_recv_chain_q
- target_recv_fused
- target_recv_response_wait
- target_response_meta_read
- target_response_received
- target_send_request
- target_spec_wait
- tree_kv_restore
- verify_accept_prep
- verify_sample_accept
- verify_setup

Hatched gaps are inferred elapsed intervals, NOT proven GPU idle. Command receive includes communication/wait. Nested spans are not additive. Raw bars retain legacy 0.02 ms display floor.

LayerSkip CodeLlama-34B + AMD-Llama-135M-code
P1 hit / p50 / UID humaneval_local_033_t0 / request 4 / step 7

exit 30 (0-based) | K1/K2 12/6 | G=M DS/PS 12/8 | C=3 | last | target GPUs=4 + draft GPU=1 | B=1, BF16, T=0.7

Measured v1.0 59f46511 | eslab18 RTX3090 | original raw-event lanes + marked gaps / next receive



1 PS roots / layout: 1.646 ms 2 Unprofiled PS dispatch: 1.191 ms (CPU gap 1.189 ms)

3 Post-PS unprofiled: 0.273 ms 4 Recorded command receive: 7.695 ms PS end -> next receive start: 7.992 ms

PS total 11.873 ms = graph 7.358 + outside graph 4.515 ms. Command -> request gap: 0.023 ms.

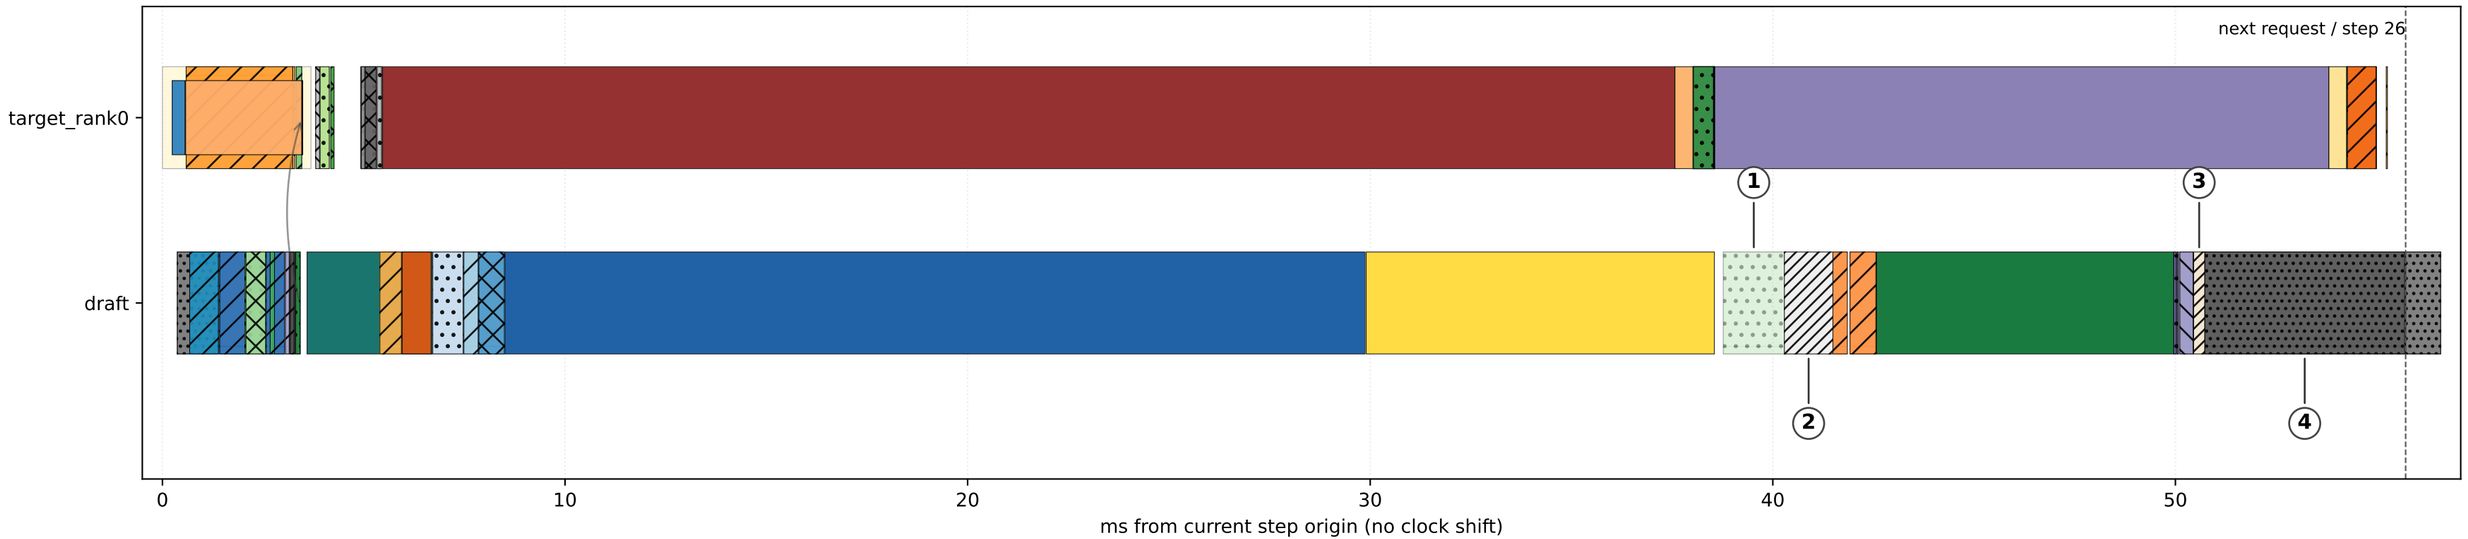
Command-to-request gap (unprofiled)	exit_logits	p1_graph_replay	proxy_send_enqueue	target_spec_wait	tree_proxy_topology_prepare
PS dispatch gap (unprofiled)	exit_proxy_launch	p1_prepare	proxy_wait	tree_accept_walk	tree_replay
Post-PS gap (unprofiled)	exit_proxy_side	p1_root_build	target_postprocess	tree_hit_pack_served_p1	tree_verify_attention_buffers
draft_recv_cmd	final_logits	p1_slot_prepare	target_recv_fused	tree_hit_parent_q_gather_p1	tree_verify_mask_prepare
draft_recv_request	glue	p2_cache_merge	target_recv_parent_q_p1	tree_hit_rerank_p1	tree_verify_meta_cpu
draft_response_pack	graph_post	p2_graph_replay	target_recv_response_wait	tree_kv_restore	tree_wire_parse_validate
draft_send_fused	graph_pre	p2_output_convert	target_response_meta_read	tree_parent_q_select	verify_accept_prep
draft_send_parent_q_p1	hit_cache_respond_hit_k1	p2_prepare	target_response_received	tree_prep	verify_sample_accept
draft_send_response	merge_cache	phase2_build	target_send_request	tree_proxy_graph_replay	verify_setup

Hatched gaps are inferred elapsed intervals, NOT proven GPU idle. Command receive includes communication/wait. Nested spans are not additive. Raw bars retain legacy 0.02 ms display floor.

LayerSkip CodeLlama-34B + AMD-Llama-135M-code
P2 hit / p50 / UID humaneval_local_023_t0 / request 3 / step 25

exit 30 (0-based) | K1/K2 12/6 | G=M DS/PS 12/8 | C=3 | last | target GPUs=4 + draft GPU=1 | B=1, BF16, T=0.7

Measured v1.0 59f46511 | eslab18 RTX3090 | original raw-event lanes + marked gaps / next receive



1 PS roots / layout: 1.520 ms 2 Unprofiled PS dispatch: 1.205 ms (CPU gap 1.205 ms)

3 Post-PS unprofiled: 0.274 ms 4 Recorded command receive: 4.969 ms PS end -> next receive start: 5.266 ms

PS total 11.905 ms = graph 7.380 + outside graph 4.525 ms. Command -> request gap: 0.022 ms.

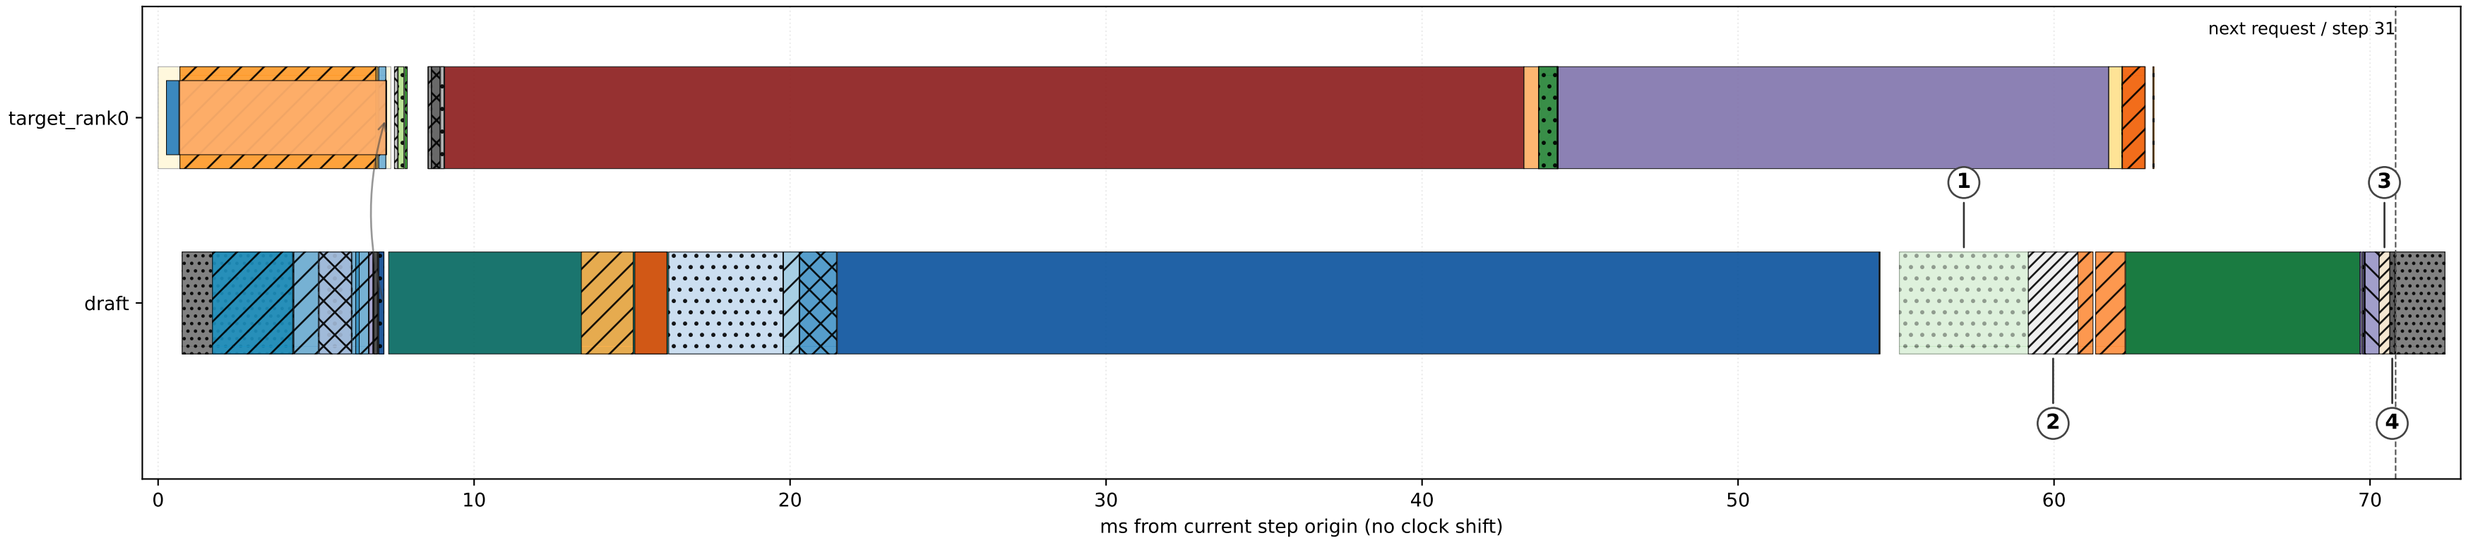
Command-to-request gap (unprofiled)	exit_logits	p1_graph_replay	proxy_send_enqueue	target_spec_wait	tree_proxy_topology_prepare
PS dispatch gap (unprofiled)	exit_proxy_launch	p1_prepare	proxy_wait	tree_accept_walk	tree_replay
Post-PS gap (unprofiled)	exit_proxy_side	p1_root_build	target_postprocess	tree_hit_pack_served_p2	tree_verify_attention_buffers
draft_recv_cmd	final_logits	p1_slot_prepare	target_recv_fused	tree_hit_parent_q_gather_p2	tree_verify_mask_prepare
draft_recv_request	glue	p2_cache_merge	target_recv_parent_q_p2	tree_hit_rerank_p2	tree_verify_meta_cpu
draft_response_pack	graph_post	p2_graph_replay	target_recv_response_wait	tree_kv_restore	tree_wire_parse_validate
draft_send_fused	graph_pre	p2_output_convert	target_response_meta_read	tree_parent_q_select	verify_accept_prep
draft_send_parent_q_p2	hit_cache_respond_hit_k2	p2_prepare	target_response_received	tree_prep	verify_sample_accept
draft_send_response	merge_cache	phase2_build	target_send_request	tree_proxy_graph_replay	verify_setup

Hatched gaps are inferred elapsed intervals, NOT proven GPU idle. Command receive includes communication/wait. Nested spans are not additive. Raw bars retain legacy 0.02 ms display floor.

CodeLlama-34B + AMD-135M-code / throughput alternative
P1 hit / p50 / UID humaneval_local_013_t0 / request 2 / step 30

exit 30 (0-based) | K1/K2 16/6 | G=M DS/PS 16/8 | C=3 | last | target GPUs=4 + draft GPU=1 | B=1, BF16, T=0.7

Measured v1.0 59f46511 | eslab18 RTX3090 | original raw-event lanes + marked gaps / next receive



1 PS roots / layout: 4.079 ms 2 Unprofiled PS dispatch: 1.571 ms (CPU gap 1.571 ms)

3 Post-PS unprofiled: 0.329 ms 4 Recorded command receive: 0.164 ms PS end -> next receive start: 0.518 ms

PS total 15.815 ms = graph 7.409 + outside graph 8.406 ms. Command -> request gap: 0.025 ms.

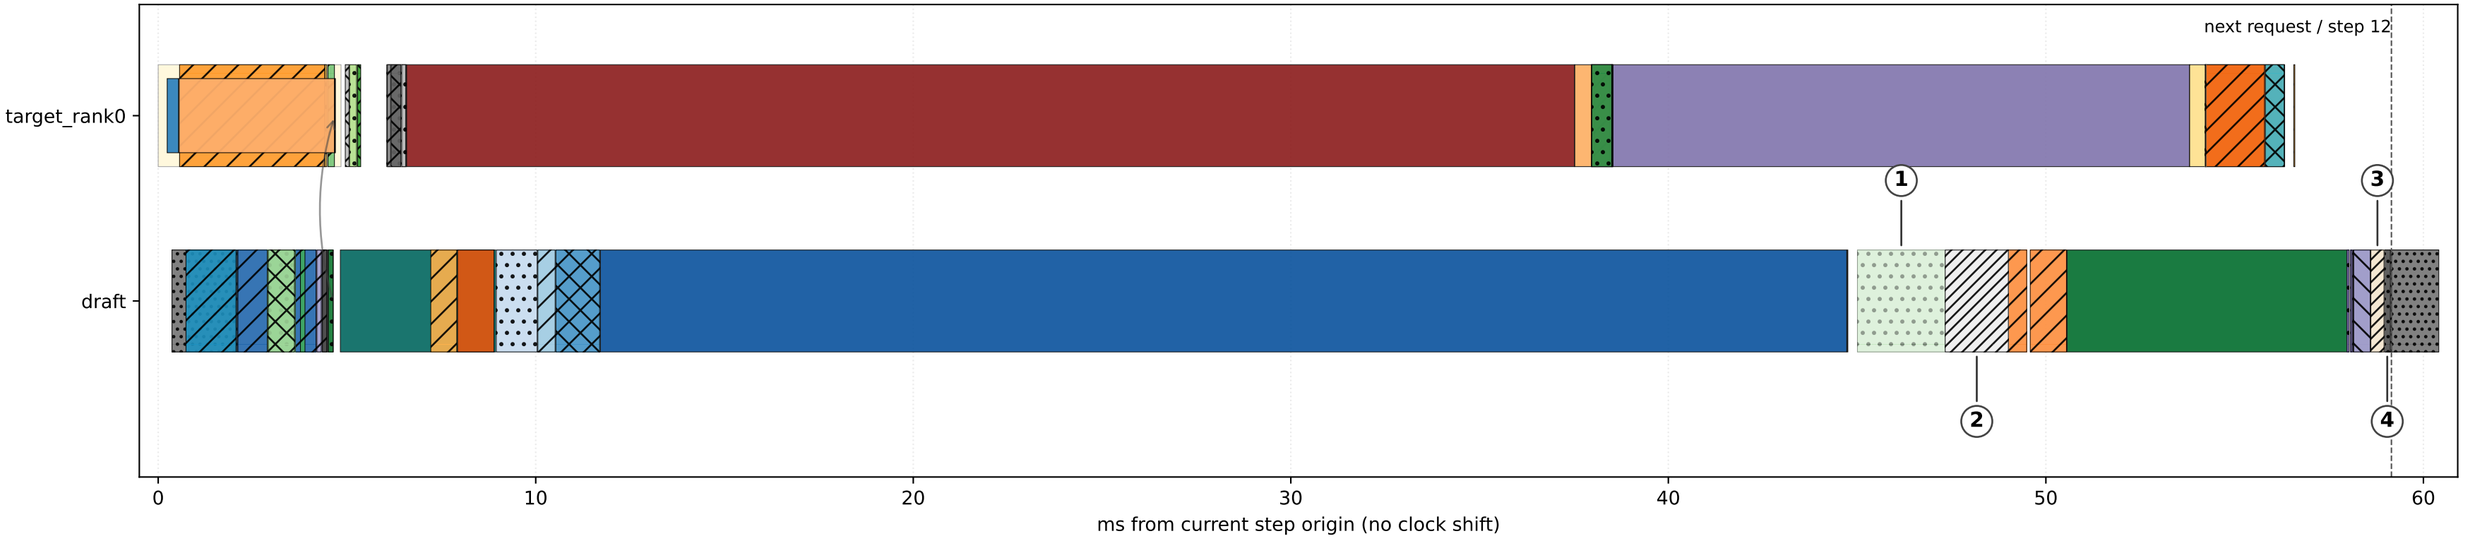
Command-to-request gap (unprofiled)	exit_logits	p1_graph_replay	proxy_send_enqueue	target_spec_wait	tree_proxy_topology_prepare
PS dispatch gap (unprofiled)	exit_proxy_launch	p1_prepare	proxy_wait	tree_accept_walk	tree_replay
Post-PS gap (unprofiled)	exit_proxy_side	p1_root_build	target_postprocess	tree_hit_pack_served_p1	tree_verify_attention_buffers
draft_recv_cmd	final_logits	p1_slot_prepare	target_recv_fused	tree_hit_parent_q_gather_p1	tree_verify_mask_prepare
draft_recv_request	glue	p2_cache_merge	target_recv_parent_q_p1	tree_hit_rerank_p1	tree_verify_meta_cpu
draft_response_pack	graph_post	p2_graph_replay	target_recv_response_wait	tree_kv_restore	tree_wire_parse_validate
draft_send_fused	graph_pre	p2_output_convert	target_response_meta_read	tree_parent_q_select	verify_accept_prep
draft_send_parent_q_p1	hit_cache_respond_hit_k1	p2_prepare	target_response_received	tree_prep	verify_sample_accept
draft_send_response	merge_cache	phase2_build	target_send_request	tree_proxy_graph_replay	verify_setup

Hatched gaps are inferred elapsed intervals, NOT proven GPU idle. Command receive includes communication/wait. Nested spans are not additive. Raw bars retain legacy 0.02 ms display floor.

CodeLlama-34B + AMD-135M-code / throughput alternative
P2 hit / p50 / UID humaneval_local_033_t0 / request 4 / step 11

exit 30 (0-based) | K1/K2 16/6 | G=M DS/PS 16/8 | C=3 | last | target GPUs=4 + draft GPU=1 | B=1, BF16, T=0.7

Measured v1.0 59f46511 | eslab18 RTX3090 | original raw-event lanes + marked gaps / next receive



1 PS roots / layout: 2.325 ms 2 Unprofiled PS dispatch: 1.673 ms (CPU gap 1.673 ms)

3 Post-PS unprofiled: 0.349 ms 4 Recorded command receive: 0.171 ms PS end -> next receive start: 0.545 ms

PS total 13.870 ms = graph 7.411 + outside graph 6.459 ms. Command -> request gap: 0.025 ms.

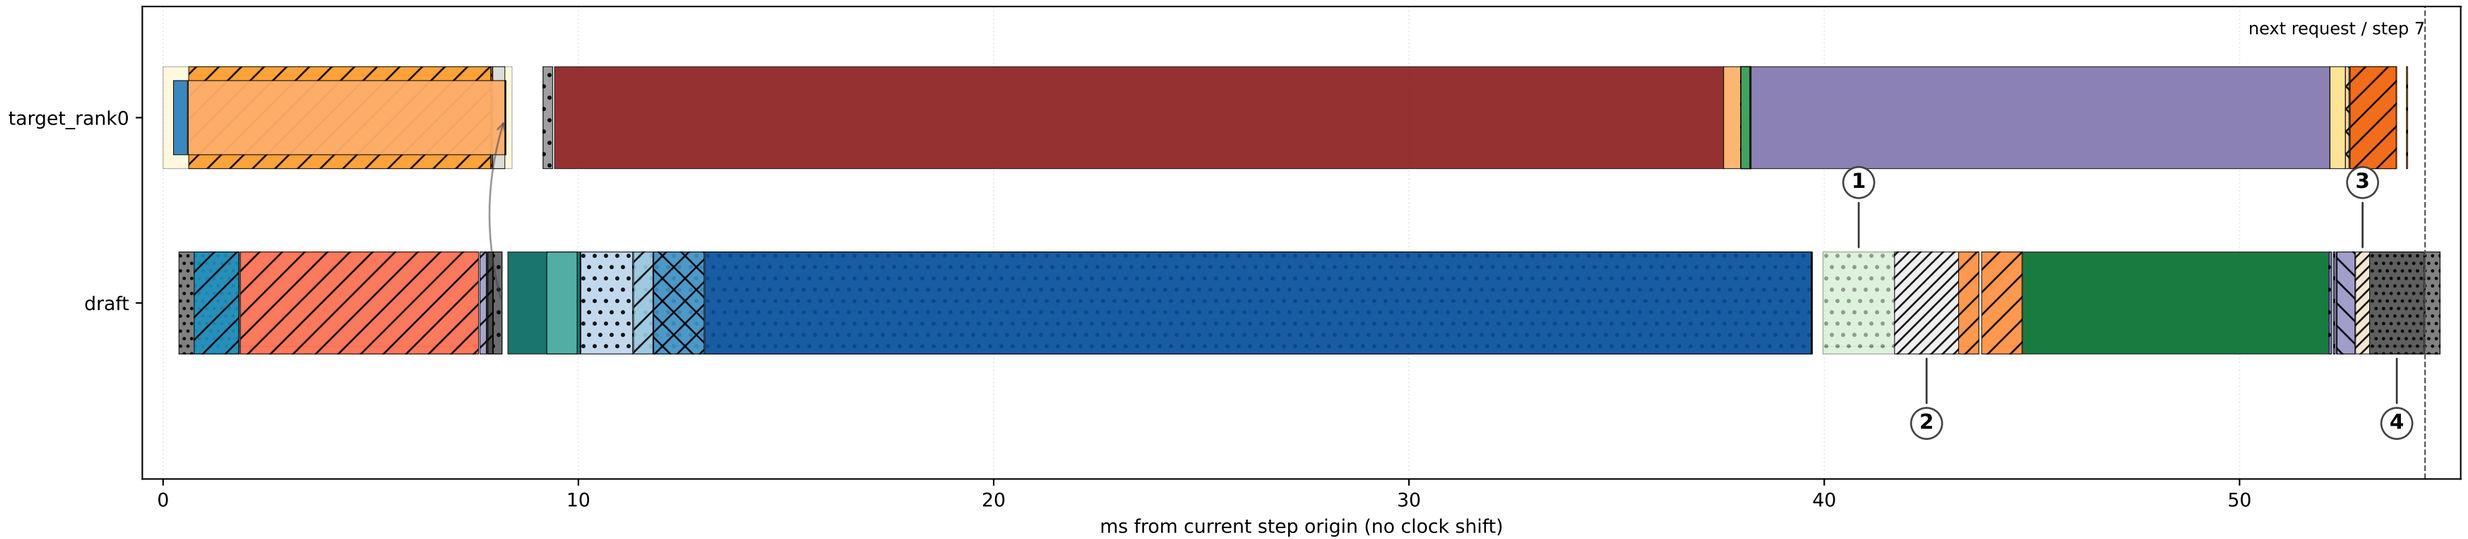
Command-to-request gap (unprofiled)
PS dispatch gap (unprofiled)
Post-PS gap (unprofiled)
draft_recv_cmd
draft_recv_request
draft_response_pack
draft_send_fused
draft_send_parent_q_p2
draft_send_response
exit_logits
exit_proxy_launch
exit_proxy_side
final_logits
glue
graph_post
graph_pre
hit_cache_respond_hit_k2
merge_cache
p1_graph_replay
p1_prepare
p1_root_build
p1_slot_prepare
p2_cache_merge
p2_graph_replay
p2_output_convert
p2_prepare
phase2_build
proxy_send_enqueue
proxy_wait
target_postprocess
target_recv_fused
target_recv_parent_q_p2
target_recv_response_wait
target_response_meta_read
target_response_received
target_send_request
target_spec_wait
tree_accept_walk
tree_hit_pack_served_p2
tree_hit_parent_q_gather_p2
tree_hit_rerank_p2
tree_kv_commit
tree_kv_restore
tree_parent_q_select
tree_prep
tree_proxy_graph_replay
tree_proxy_topology_prepare
tree_replay
tree_verify_attention_buffers
tree_verify_mask_prepare
tree_verify_meta_cpu
tree_wire_parse_validate
verify_accept_prep
verify_sample_accept
verify_setup

Hatched gaps are inferred elapsed intervals, NOT proven GPU idle. Command receive includes communication/wait. Nested spans are not additive. Raw bars retain legacy 0.02 ms display floor.

CodeLlama-34B + AMD-135M-code / throughput alternative
Miss / p50 / UID humaneval_local_023_t0 / request 3 / step 6

exit 30 (0-based) | K1/K2 16/6 | G=M DS/PS 16/8 | C=3 | last | target GPUs=4 + draft GPU=1 | B=1, BF16, T=0.7

Measured v1.0 59f46511 | eslab18 RTX3090 | original raw-event lanes + marked gaps / next receive



1 PS roots / layout: 1.726 ms 2 Unprofiled PS dispatch: 1.542 ms (CPU gap 1.544 ms)

3 Post-PS unprofiled: 0.336 ms 4 Recorded command receive: 1.312 ms PS end -> next receive start: 1.672 ms

PS total 13.096 ms = graph 7.375 + outside graph 5.721 ms. Command -> request gap: 0.024 ms.

- Command-to-request gap (unprofiled)
- PS dispatch gap (unprofiled)
- Post-PS gap (unprofiled)
- chain_accept
- chain_proxy_graph_replay
- draft_glue_replay
- draft_recv_cmd
- draft_recv_request
- draft_response_pack
- draft_send_chain_q
- draft_send_fused
- draft_send_response
- exit_logits
- exit_proxy_launch
- exit_proxy_side
- final_logits
- glue
- graph_post
- graph_pre
- hit_cache_respond_miss
- merge_cache
- p1_graph_replay
- p1_prepare
- p1_root_build
- p1_slot_prepare
- p2_cache_merge
- p2_graph_replay
- p2_output_convert
- p2_prepare
- phase1_build
- phase2_build
- proxy_send_enqueue
- proxy_wait
- target_postprocess
- target_recv_chain_q
- target_recv_fused
- target_recv_response_wait
- target_response_meta_read
- target_response_received
- target_send_request
- target_spec_wait
- tree_kv_restore
- verify_accept_prep
- verify_sample_accept
- verify_setup

Hatched gaps are inferred elapsed intervals, NOT proven GPU idle. Command receive includes communication/wait. Nested spans are not additive. Raw bars retain legacy 0.02 ms display floor.