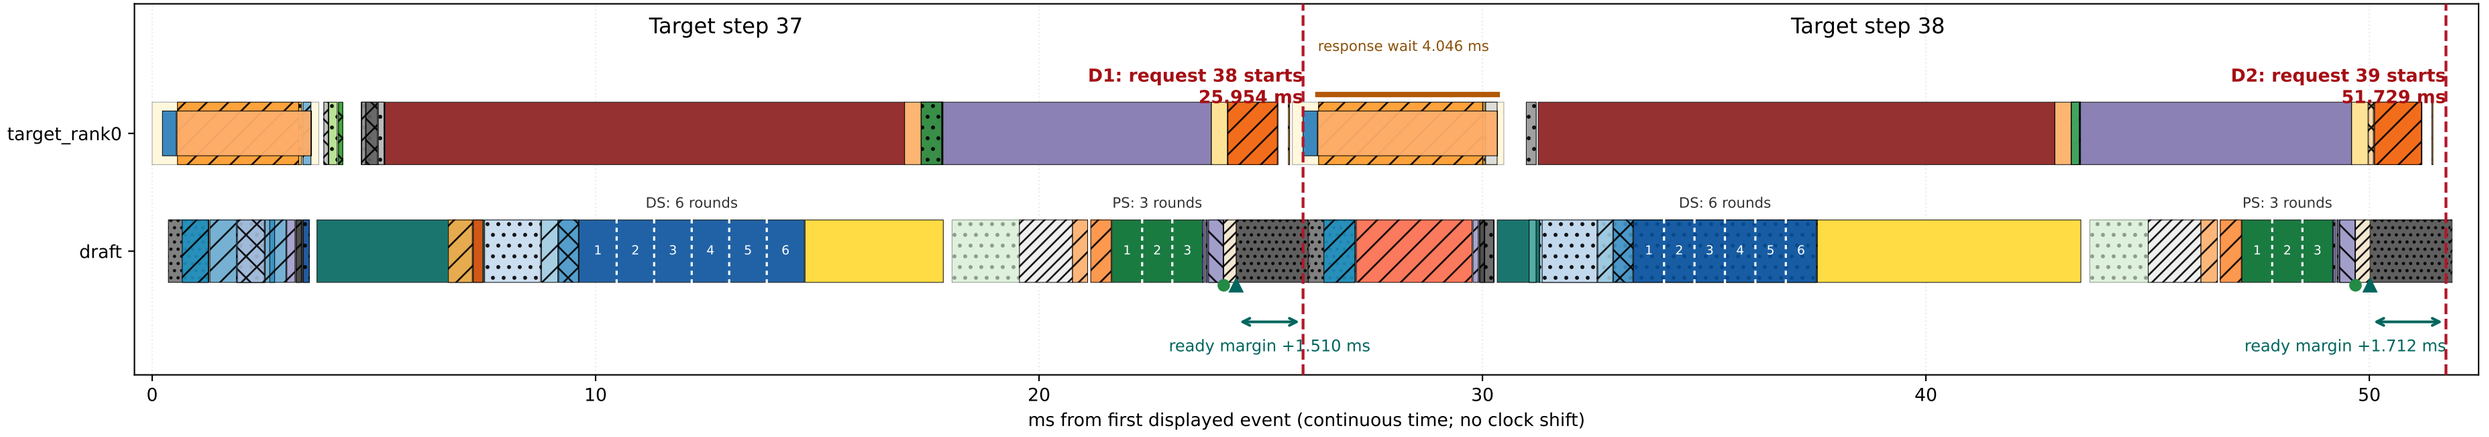
LayerSkip Llama2-7B + llama-68M | Selected setting | P1 hit | UID 288_t0 | steps 37-38

Historical v1.0 82d964e0 | exit 20 (0-based) | K1/K2 6/3 | G=M DS/PS 12/8 | C=3 | frontier | target GPUs=1 + draft GPU=1 | B=1, BF16, T=0.7

Red dashed line = next target request starts (observed deadline reference). Green circle = PS cache ready; triangle = draft ready to receive.

Both target steps and their DS/PS work share one time axis. Only the following request boundary is used for the second deadline.

In this measurement: no observed proxy computation / target graph_post overlap. Send enqueue does not measure transfer completion.



ALL ELIGIBLE STEPS FOR THIS SETTING (not the two example steps above); Overall is step-weighted.

Current outcome	Steps	Rate (%)	AL (tokens/step)	Profile TPS (tokens/s)
P1 hit	115	35.826	2.817	109.397
P2 hit	164	51.090	2.366	93.922
Miss	42	13.084	2.024	79.867
Overall	**321**	**86.916**	**2.483**	**97.707**

Rate = source / all steps; Overall rate = total hit (P1 + P2) / all. AL includes recovery/bonus. Profile TPS = emitted tokens / summed step seconds.

Current outcome	Steps	Step (ms)	DS total (ms)	DS spare* (ms)	DS graph/K1 (ms)	PS total (ms)	PS spare (ms)	PS graph/K2 (ms)
P1 hit	115	25.754	7.104	+3.588	0.830	6.378	+1.638	0.679
P2 hit	164	25.190	6.939	+4.204	0.802	6.428	+1.369	0.681
Miss	42	25.340	6.327	+5.751	0.689	6.211	+1.687	0.680
Overall	**321**	**25.411**	**6.918**	**+4.186**	**0.797**	**6.382**	**+1.507**	**0.680**

Step = current target request start -> next request start. *DS spare = proxy send-ready - DS end; PS spare = next request start - draft receive-ready.

Positive spare = early; negative = late. Graph/K includes sampling/selection, not isolated transformer latency. These are profiled-subset TPS, not full evaluation TPS.

Command-to-request gap (unprofiled)
PS dispatch gap (unprofiled)
PS prepare A: roots / visibility
PS prepare B: graph input buffers
Post-PS gap (unprofiled)
chain_accept
chain_proxy_graph_replay
draft_glue_replay
draft_recv_cmd
draft_recv_request
draft_response_pack
draft_send_chain_q
draft_send_fused
draft_send_parent_q_p1
draft_send_response
exit_logits
exit_proxy_launch
exit_proxy_side
final_logits
glue
graph_post
graph_pre
hit_cache_respond_hit_k1
hit_cache_respond_miss
merge_cache
p1_graph_replay
p1_prepare
p1_root_build
p1_slot_prepare
p2_cache_merge
p2_graph_replay
p2_output_convert
phase1_build
phase2_build
proxy_send_enqueue
proxy_wait
target_postprocess
target_recv_chain_q
target_recv_fused
target_recv_parent_q_p1
target_recv_response_wait
target_response_meta_read
target_response_received
target_send_request
target_spec_wait
tree_accept_walk
tree_hit_pack_served_p1
tree_hit_parent_q_gather_p1
tree_hit_rerank_p1
tree_kv_restore
tree_parent_q_select
tree_prep
tree_proxy_graph_replay
tree_proxy_topology_prepare
tree_replay
tree_verify_attention_buffers
tree_verify_mask_prepare
tree_verify_meta_cpu
tree_wire_parse_validate
verify_accept_prep
verify_sample_accept
verify_setup

Hatched gaps = unprofiled elapsed intervals, NOT proven GPU idle. Nested spans are not additive. White graph dividers show counts, NOT per-forward measurements.

Raw intervals retained, no 0.02 ms visibility floor. Cross-device margins inherit CUDA-to-host clock alignment uncertainty; exact raw events and source hashes accompany this figure.

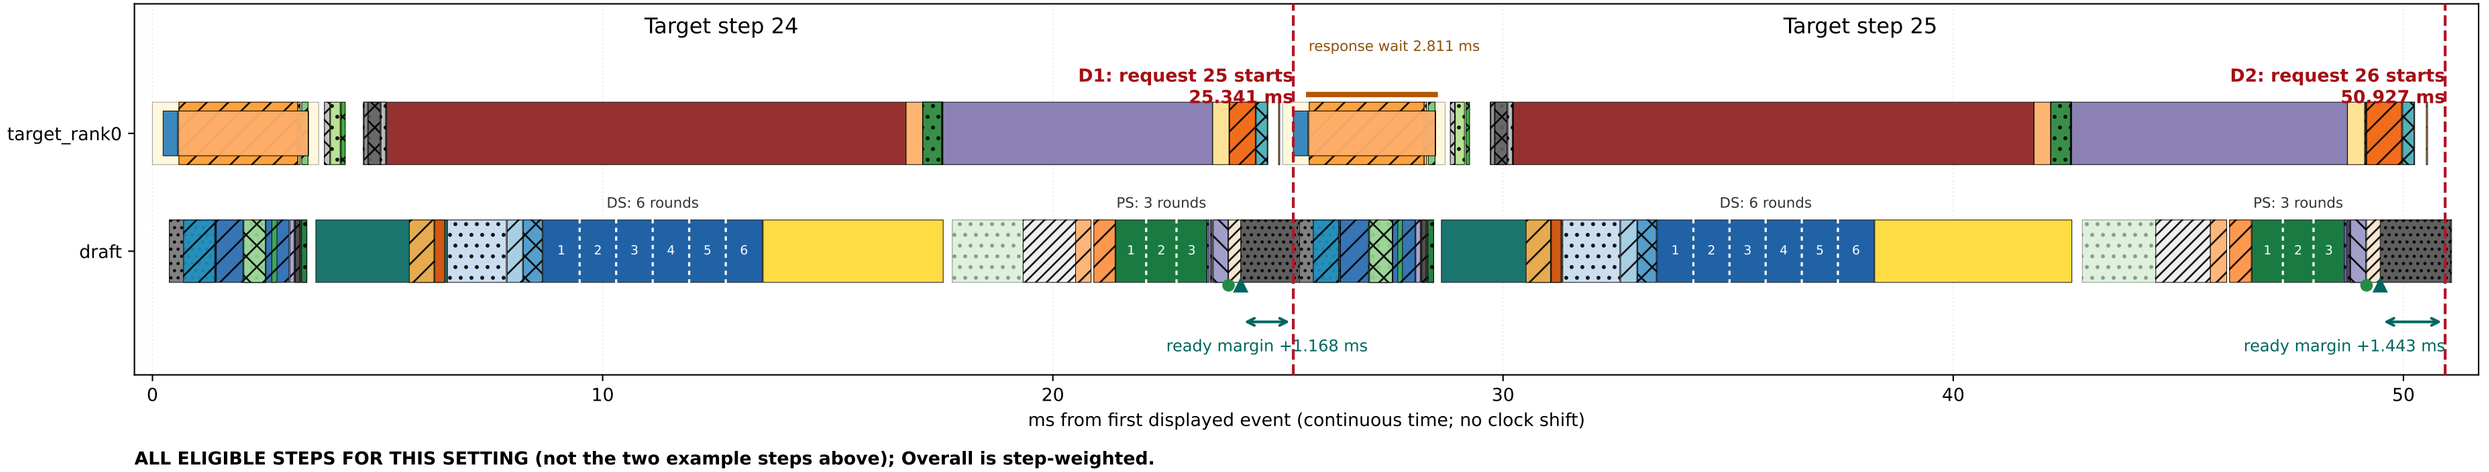
LayerSkip Llama2-7B + llama-68M | Selected setting | P2 hit | UID 244_t0 | steps 24-25

Historical v1.0 82d964e0 | exit 20 (0-based) | K1/K2 6/3 | G=M DS/PS 12/8 | C=3 | frontier | target GPUs=1 + draft GPU=1 | B=1, BF16, T=0.7

Red dashed line = next target request starts (observed deadline reference). Green circle = PS cache ready; triangle = draft ready to receive.

Both target steps and their DS/PS work share one time axis. Only the following request boundary is used for the second deadline.

In this measurement: no observed proxy computation / target graph_post overlap. Send enqueue does not measure transfer completion.



ALL ELIGIBLE STEPS FOR THIS SETTING (not the two example steps above); Overall is step-weighted.

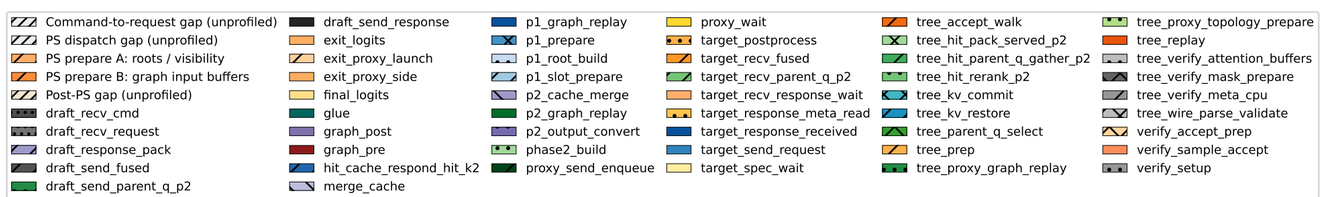
Current outcome	Steps	Rate (%)	AL (tokens/step)	Profile TPS (tokens/s)
P1 hit	115	35.826	2.817	109.397
P2 hit	164	51.090	2.366	93.922
Miss	42	13.084	2.024	79.867
Overall	**321**	**86.916**	**2.483**	**97.707**

Rate = source / all steps; Overall rate = total hit (P1 + P2) / all. AL includes recovery/bonus. Profile TPS = emitted tokens / summed step seconds.

Current outcome	Steps	Step (ms)	DS total (ms)	DS spare* (ms)	DS graph/K1 (ms)	PS total (ms)	PS spare (ms)	PS graph/K2 (ms)
P1 hit	115	25.754	7.104	+3.588	0.830	6.378	+1.638	0.679
P2 hit	164	25.190	6.939	+4.204	0.802	6.428	+1.369	0.681
Miss	42	25.340	6.327	+5.751	0.689	6.211	+1.687	0.680
Overall	**321**	**25.411**	**6.918**	**+4.186**	**0.797**	**6.382**	**+1.507**	**0.680**

Step = current target request start -> next request start. *DS spare = proxy send-ready - DS end; PS spare = next request start - draft receive-ready.

Positive spare = early; negative = late. Graph/K includes sampling/selection, not isolated transformer latency. These are profiled-subset TPS, not full evaluation TPS.



Hatched gaps = unprofiled elapsed intervals, NOT proven GPU idle. Nested spans are not additive. White graph dividers show counts, NOT per-forward measurements.

Raw intervals retained, no 0.02 ms visibility floor. Cross-device margins inherit CUDA-to-host clock alignment uncertainty; exact raw events and source hashes accompany this figure.

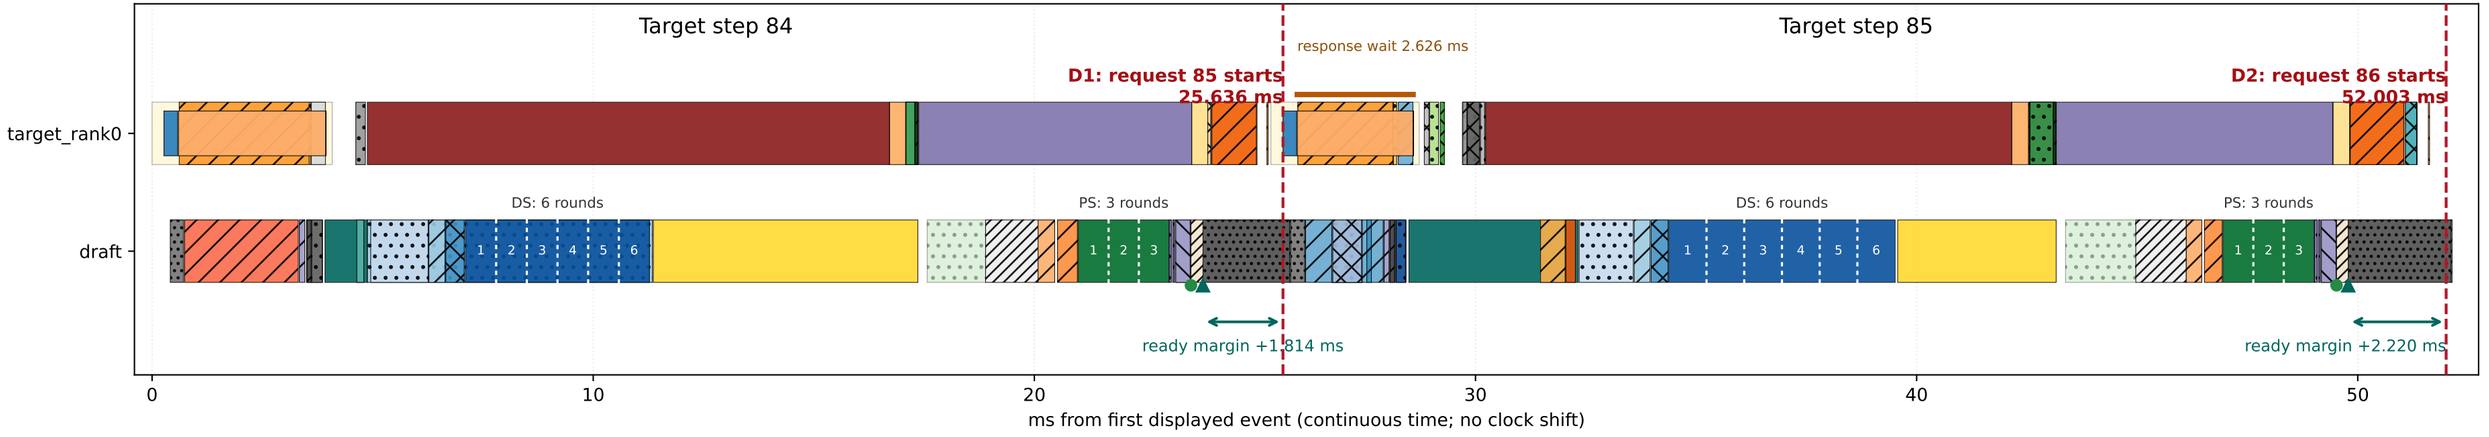
LayerSkip Llama2-7B + llama-68M | Selected setting | Miss | UID 288_t0 | steps 84-85

Historical v1.0 82d964e0 | exit 20 (0-based) | K1/K2 6/3 | G=M DS/PS 12/8 | C=3 | frontier | target GPUs=1 + draft GPU=1 | B=1, BF16, T=0.7

Red dashed line = next target request starts (observed deadline reference). Green circle = PS cache ready; triangle = draft ready to receive.

Both target steps and their DS/PS work share one time axis. Only the following request boundary is used for the second deadline.

In this measurement: no observed proxy computation / target graph_post overlap. Send enqueue does not measure transfer completion.



ALL ELIGIBLE STEPS FOR THIS SETTING (not the two example steps above); Overall is step-weighted.

Current outcome	Steps	Rate (%)	AL (tokens/step)	Profile TPS (tokens/s)
P1 hit	115	35.826	2.817	109.397
P2 hit	164	51.090	2.366	93.922
Miss	42	13.084	2.024	79.867
Overall	**321**	**86.916**	**2.483**	**97.707**

Rate = source / all steps; Overall rate = total hit (P1 + P2) / all. AL includes recovery/bonus. Profile TPS = emitted tokens / summed step seconds.

Current outcome	Steps	Step (ms)	DS total (ms)	DS spare* (ms)	DS graph/K1 (ms)	PS total (ms)	PS spare (ms)	PS graph/K2 (ms)
P1 hit	115	25.754	7.104	+3.588	0.830	6.378	+1.638	0.679
P2 hit	164	25.190	6.939	+4.204	0.802	6.428	+1.369	0.681
Miss	42	25.340	6.327	+5.751	0.689	6.211	+1.687	0.680
Overall	**321**	**25.411**	**6.918**	**+4.186**	**0.797**	**6.382**	**+1.507**	**0.680**

Step = current target request start -> next request start. *DS spare = proxy send-ready - DS end; PS spare = next request start - draft receive-ready.

Positive spare = early; negative = late. Graph/K includes sampling/selection, not isolated transformer latency. These are profiled-subset TPS, not full evaluation TPS.

Command-to-request gap (unprofiled)
PS dispatch gap (unprofiled)
PS prepare A: roots / visibility
PS prepare B: graph input buffers
Post-PS gap (unprofiled)
chain_accept
chain_proxy_graph_replay
draft_glue_replay
draft_recv_cmd
draft_recv_request
draft_response_pack
draft_send_chain_q
draft_send_fused
draft_send_parent_q_p1
draft_send_response
exit_logits
exit_proxy_launch
exit_proxy_side
final_logits
glue
graph_post
graph_pre
hit_cache_respond_hit_k1
hit_cache_respond_miss
merge_cache
p1_graph_replay
p1_prepare
p1_root_build
p1_slot_prepare
p2_cache_merge
p2_graph_replay
p2_output_convert
phase1_build
phase2_build
proxy_send_enqueue
proxy_wait
target_postprocess
target_recv_chain_q
target_recv_fused
target_recv_parent_q_p1
target_recv_response_wait
target_response_meta_read
target_response_received
target_send_request
target_spec_wait
tree_accept_walk
tree_hit_pack_served_p1
tree_hit_parent_q_gather_p1
tree_hit_rerank_p1
tree_kv_commit
tree_parent_q_select
tree_prep
tree_proxy_graph_replay
tree_proxy_topology_prepare
tree_replay
tree_verify_attention_buffers
tree_verify_mask_prepare
tree_verify_meta_cpu
tree_wire_parse_validate
verify_accept_prep
verify_sample_accept
verify_setup

Hatched gaps = unprofiled elapsed intervals, NOT proven GPU idle. Nested spans are not additive. White graph dividers show counts, NOT per-forward measurements.

Raw intervals retained, no 0.02 ms visibility floor. Cross-device margins inherit CUDA-to-host clock alignment uncertainty; exact raw events and source hashes accompany this figure.