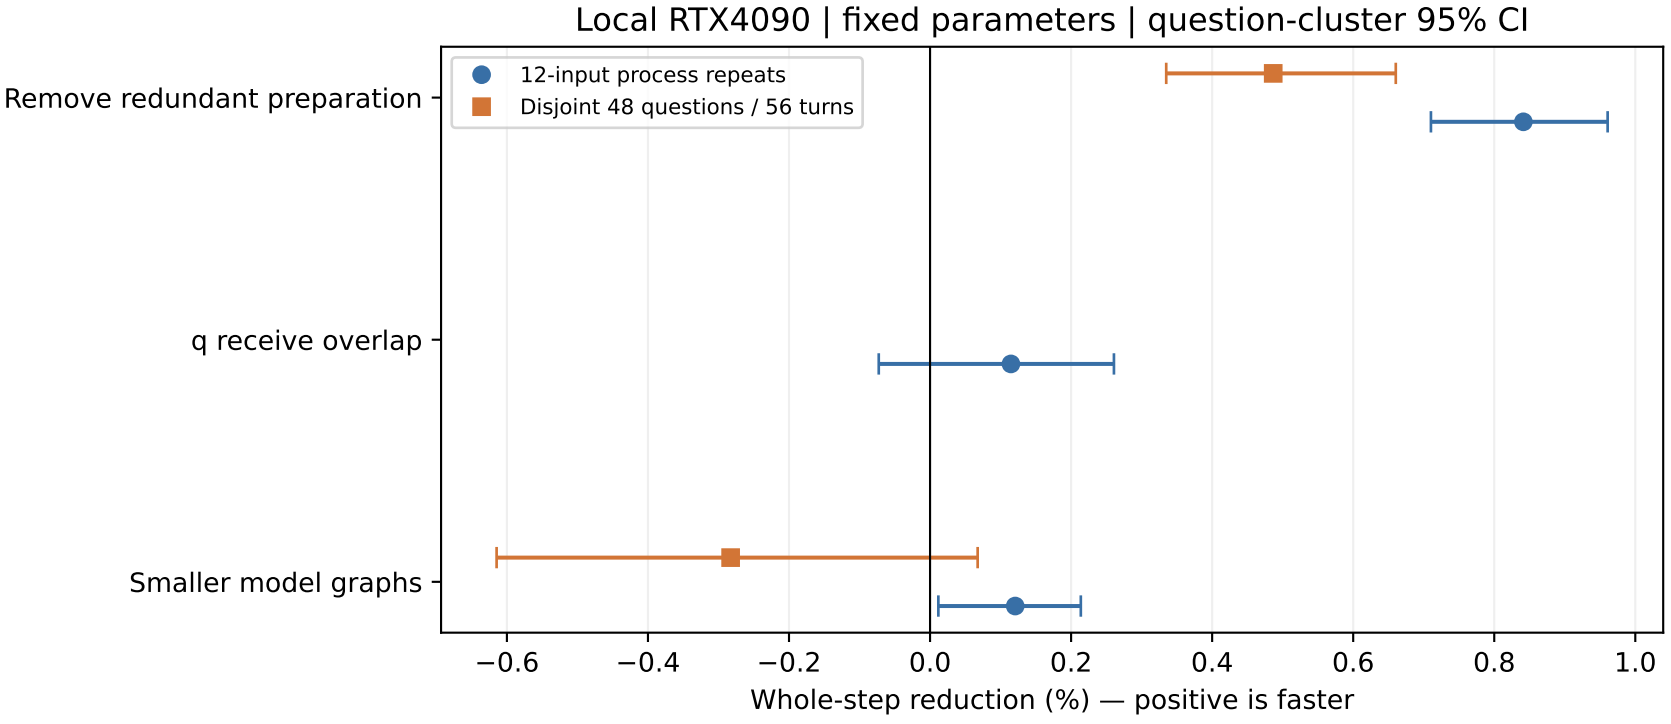
Local RTX4090 | fixed parameters | question-cluster 95% CI
12-input process repeats
Disjoint 48 questions / 56 turns
Remove redundant preparation
q receive overlap
Smaller model graphs
−0.6
−0.4
−0.2
0.0
0.2
0.4
0.6
0.8
1.0
Whole-step reduction (%) — positive is faster